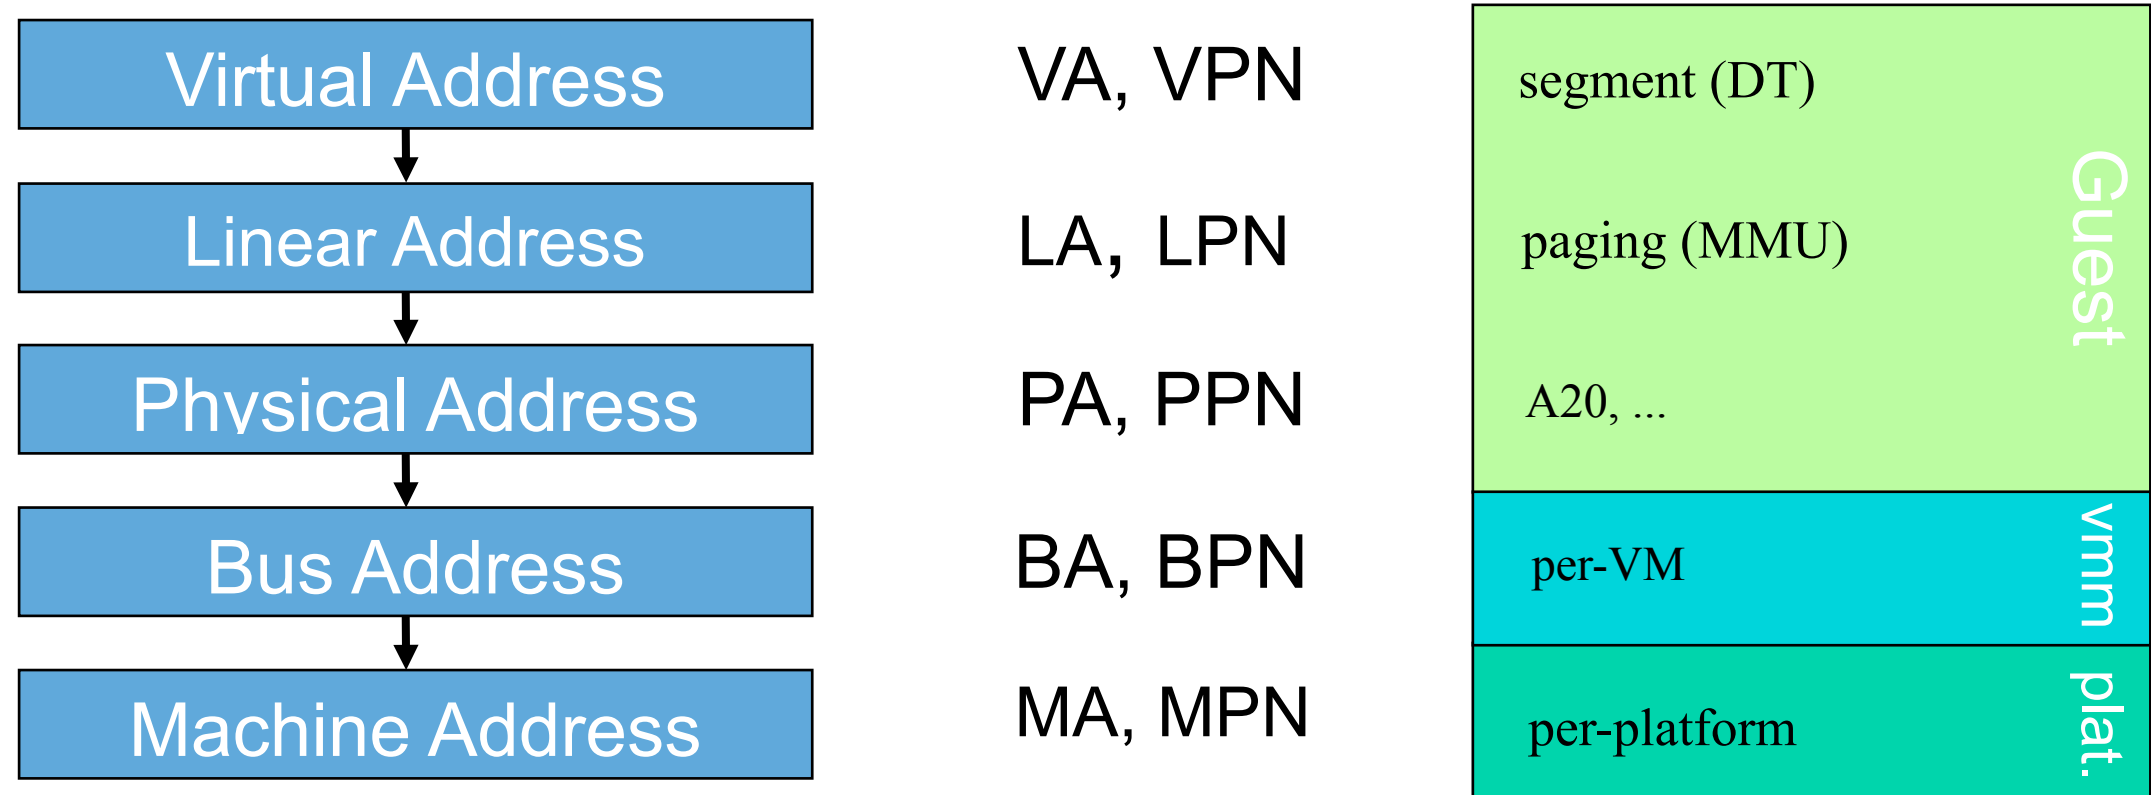


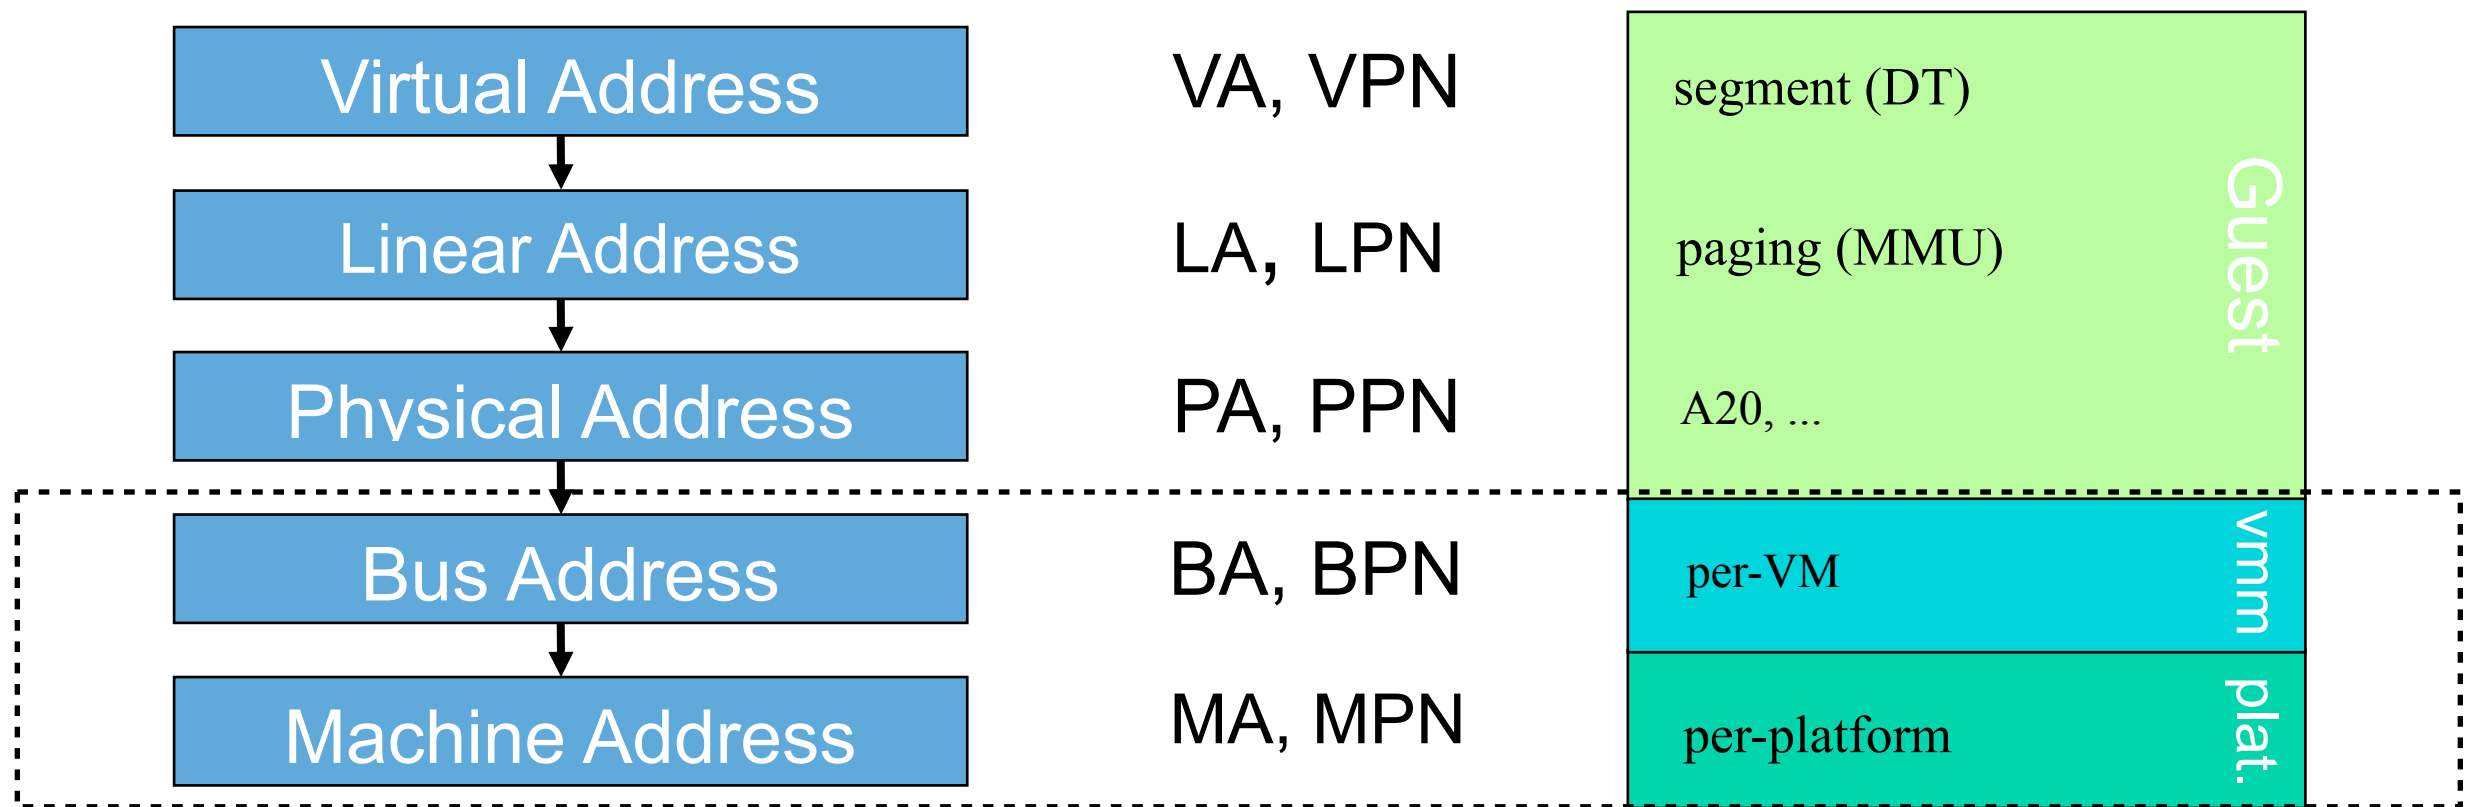
BusMem and Memory Scheduling

Alex Garthwaite
MonitorU
July 27th, 2010

Background

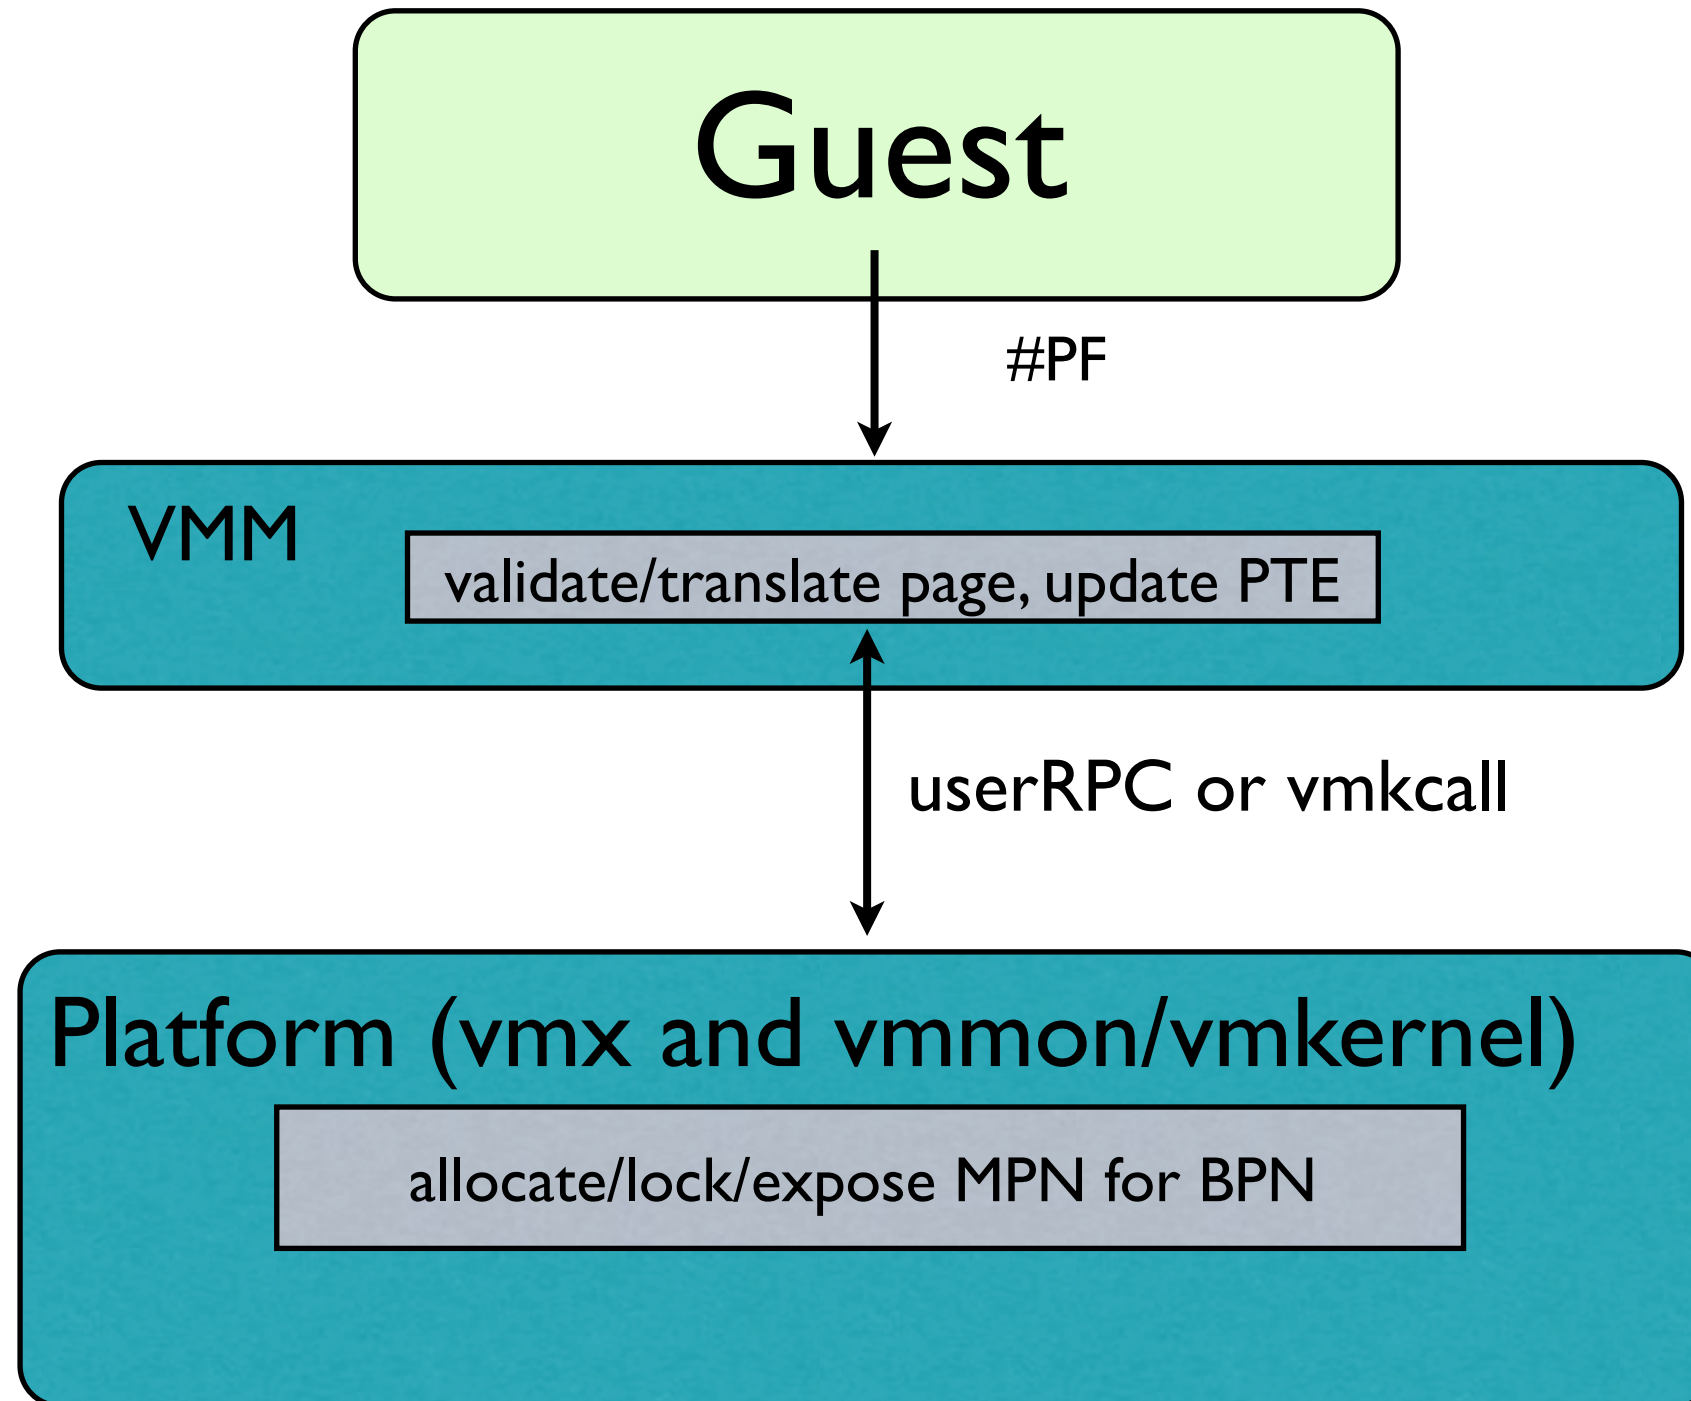


Background



BusMem: map BPNs to MPNs and track state

Guest Memory



Topics

- Memory state in the vmm
- Translation/zapping/invalidation
- Memory scheduling and services
- Some new directions

Not Covered


- VM states and allocation
 - power-on/off
 - checkpointing (of memory)/migration
- How traces on memory are used
- Overheadmem and VM admission
- Monitor actions
- Details of large/small page management
- Anon page allocation/release
- ...

Themes

- Scaling and BusMem lock contention
- Overheads from tracking BPN state
- Release of memory and monitor actions
 - swapping/page-sharing/breaking sharing
- Coordination of policy/mechanisms across layers in and across VMs and platforms

BusMem Frames

BusMemFrame



BusMem Frames

BusMemFrame


MPN
flags
props/traces

(mostly) caching
platform state


dirty
dirtyLarge
backedLarge
mappedLarge
ballooned
sampled
sampledTouched
sampledDirty
sampledBalloon
breakSharing

BusMem Frames

BusMemFrame



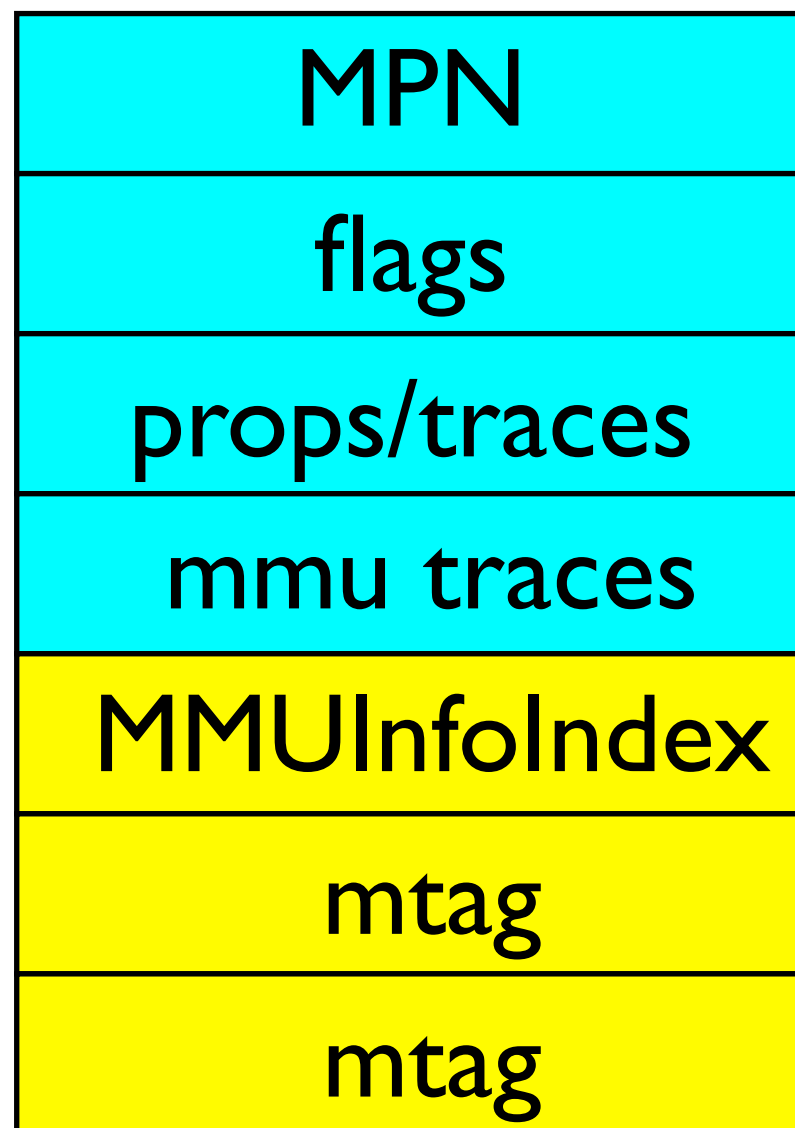
BusMemTrace



BusMem Frames

BusMemFrame

sw mmu



vcpuset

⋮

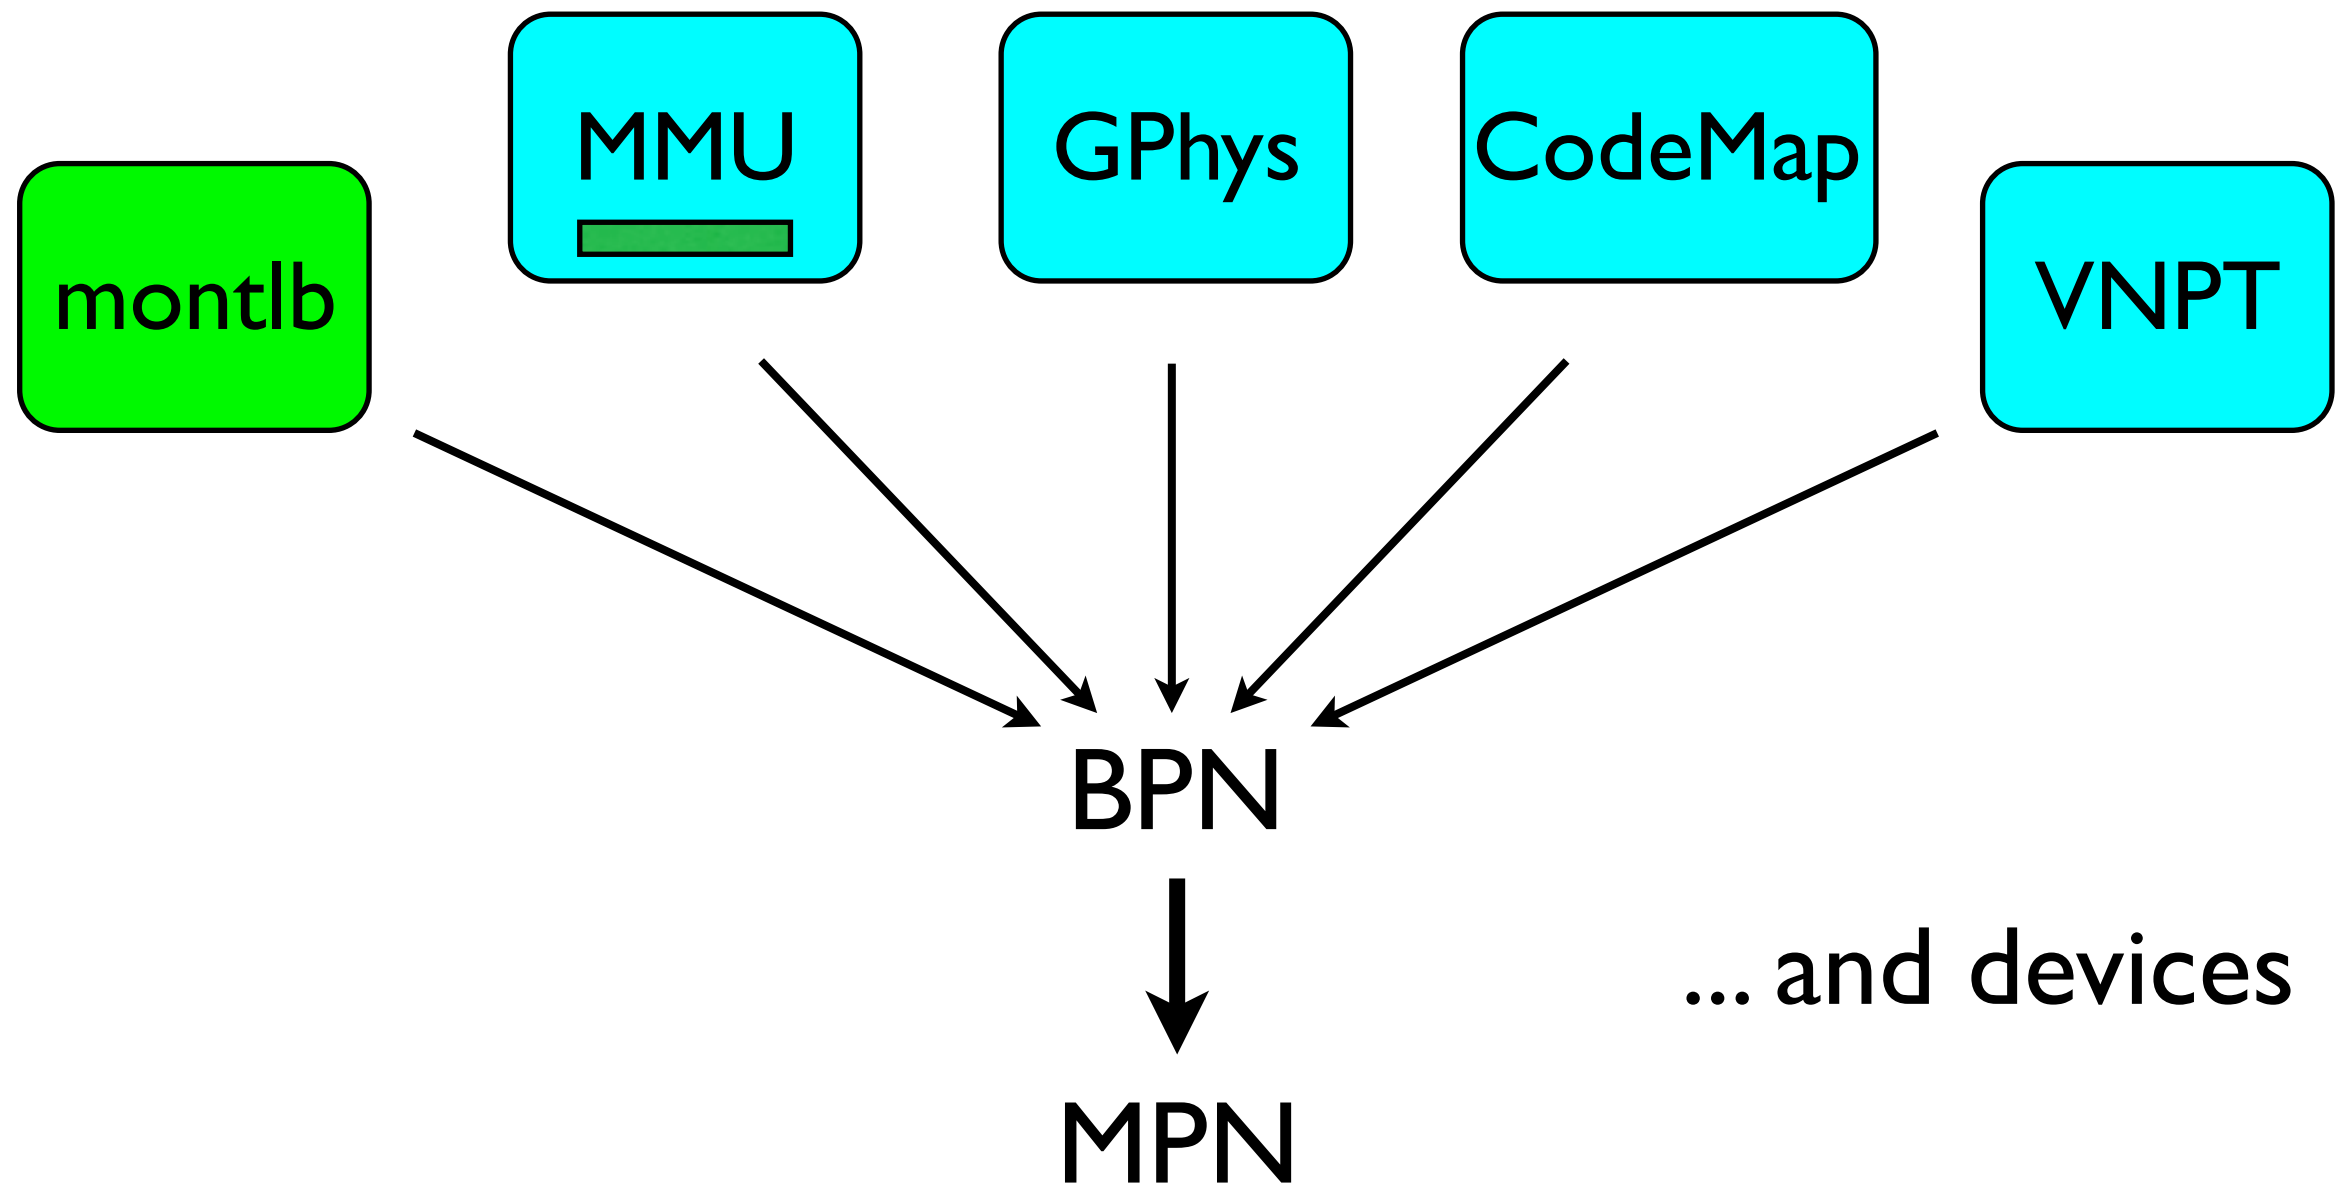
Translation/Validation

- Translation: BPN => MPN from BusMem
- Validation: call platform for BPN => MPN
 read MPN
 if (slow-path check) {
 do {
 open transaction if MPN invalid, breakcow
 Platform_LockGuestPage to get MPN if
 invalid, breakcow, or needs dirtying
 BusMemLock();
 validate MPN/state of BPN
 BusMemUnlock();
 } while (MPN not valid);
 }
 pin BPN

Translation/Validation

- Slow-path check
 - INVALID_MPN, breakcow and shared, sampling, needs dirtying
- Breaking sharing
 - frame->breakSharing = 1
 - take BusMemLock, invalidate, clear MPN/COW, release BusMemLock
- Validation checks

BusMem Clients



Locking/Pinning

- Locking: fix mapping (BPN => MPN)
- Pinning: prevent release of mapping
 - examples: ESX p2mCache, vmx PhysMem, BusMem

PinCounts (shared)

BPN	count
...	...

PinSlots (per-vcpu)

BPN	
...	(montlb)
	(mmu)

Invalidation/Zapping

- Invalidation: request zap, allow for MPN to change or be released
- Zapping: notify clients to drop cached information about BPNs
- Protocol to support translation fast-path
 - `MPN <= INVALID_MPN` before crosscall
- Coalescing concurrent invalidations and crosscall/lock contention

Invalidation/Zapping

- Five kinds of invalidation
 - PageList from {swapping, page-sharing, p2mUpdate, remap}
 - range of BPNs from {large pages, p2mUpdate, region-invalidation}
 - single BPN from {balloon, breakcow}
 - sampling
 - physMem map/unmap

Invalidation/Zapping

- Steps:
 - acquire BusMem lock
 - mark BPNs' frames (INVALID_MPN)
 - check for pinned pages/mark cases
 - zap unpinned pages on all vcpus
 - clean up BPNs' frames
 - release BusMem lock

MemSched Services

release

page-sharing

ballooning

swapping

reallocation

remapping

p2mUpdate

VMCI

usage

sampling

checking

Page-sharing

- 100ms periodic task
 - randomly sample, hash, share collisions
 - platform match hints
- Cost of breaking sharing
 - p2mUpdate mechanism
- Invariants:
 - stop so frameMPN usable
 - rate-limit to minimize impact on vcpu
- Java ballooning and guest-hints

Ballooning

guest knowledge

asynchronous/can be slow

reserves PPNs

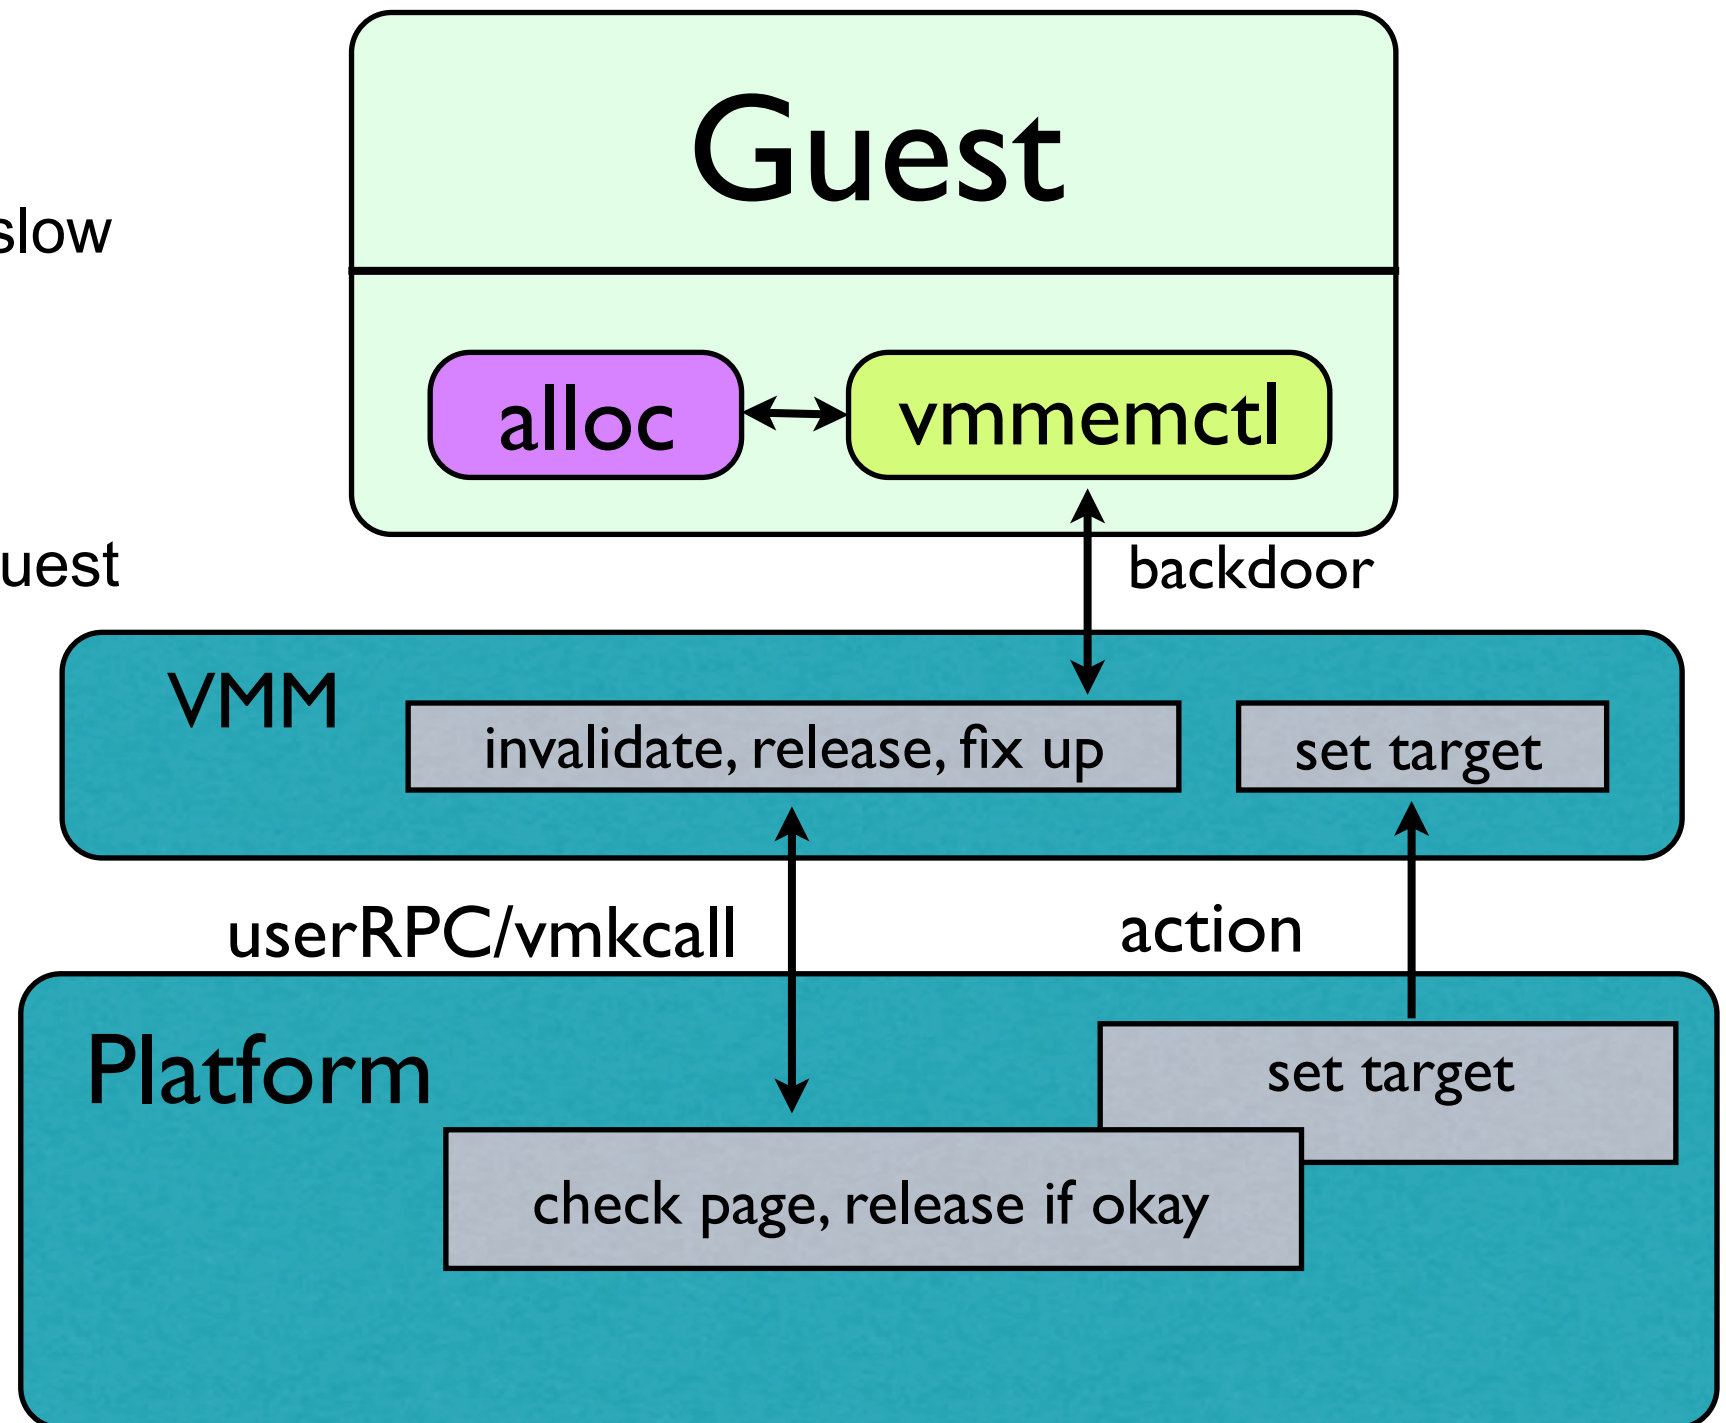
PPNs determined by guest

may not reclaim MPNs

may be reset by guest

relatively inflexible

shared zeroPage



Swapping

random choice

synchronous/transparent

one userRPC per swap-set

guest touches page in memory

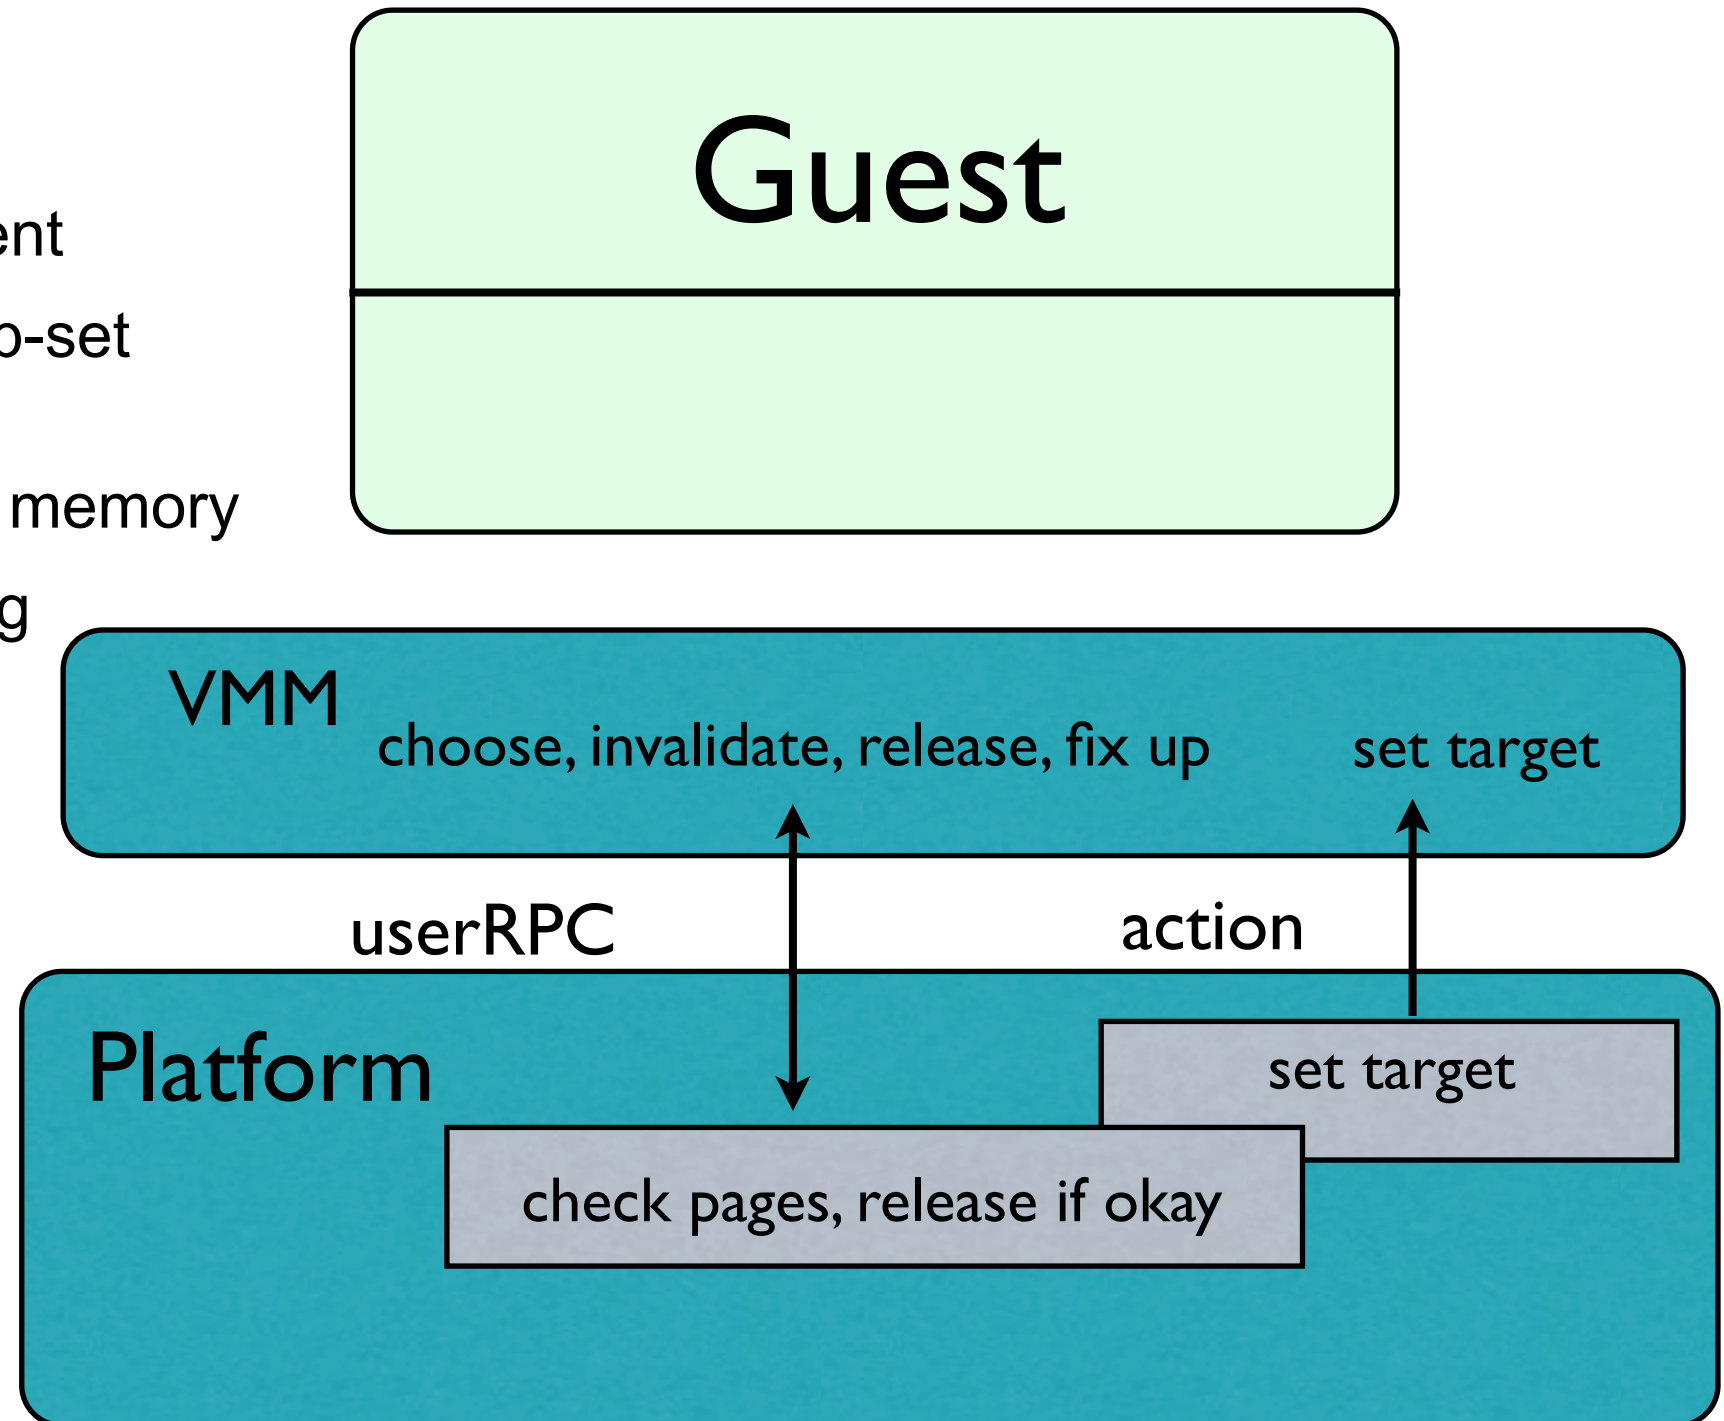
interaction w/ballooning

MN changes:

selection in platform

swap/allocation rates

avoid stopping vcpus



Remapping


- Use cases:
 - NUMA-migration
 - large-page defragmentation
 - page-retirement
 - page-coloring
- Get PageList batch, invalidate, call platform for new MPNs, update BPNs' frames

Sampling


- Four sets of 100 pages each, randomly
- Once per minute, oldest set killed, new one created
- Tracks avg. working-set size
 - used to balance memory across VMs
- Zaps without checking pinned state

MemSched Services


ballooning
(from guest)



sharing
(10Hz)



swapping
(swap action)




MemSched

Cooperative model with platform scheduler

- > declarative model (currently)
- > balance memory across VMs
 - based on memory usage (active wss, idle-page tax)
 - approx. lottery scheduling
 - once per second

`memschedInfo` in
`sharedArea`
and actions



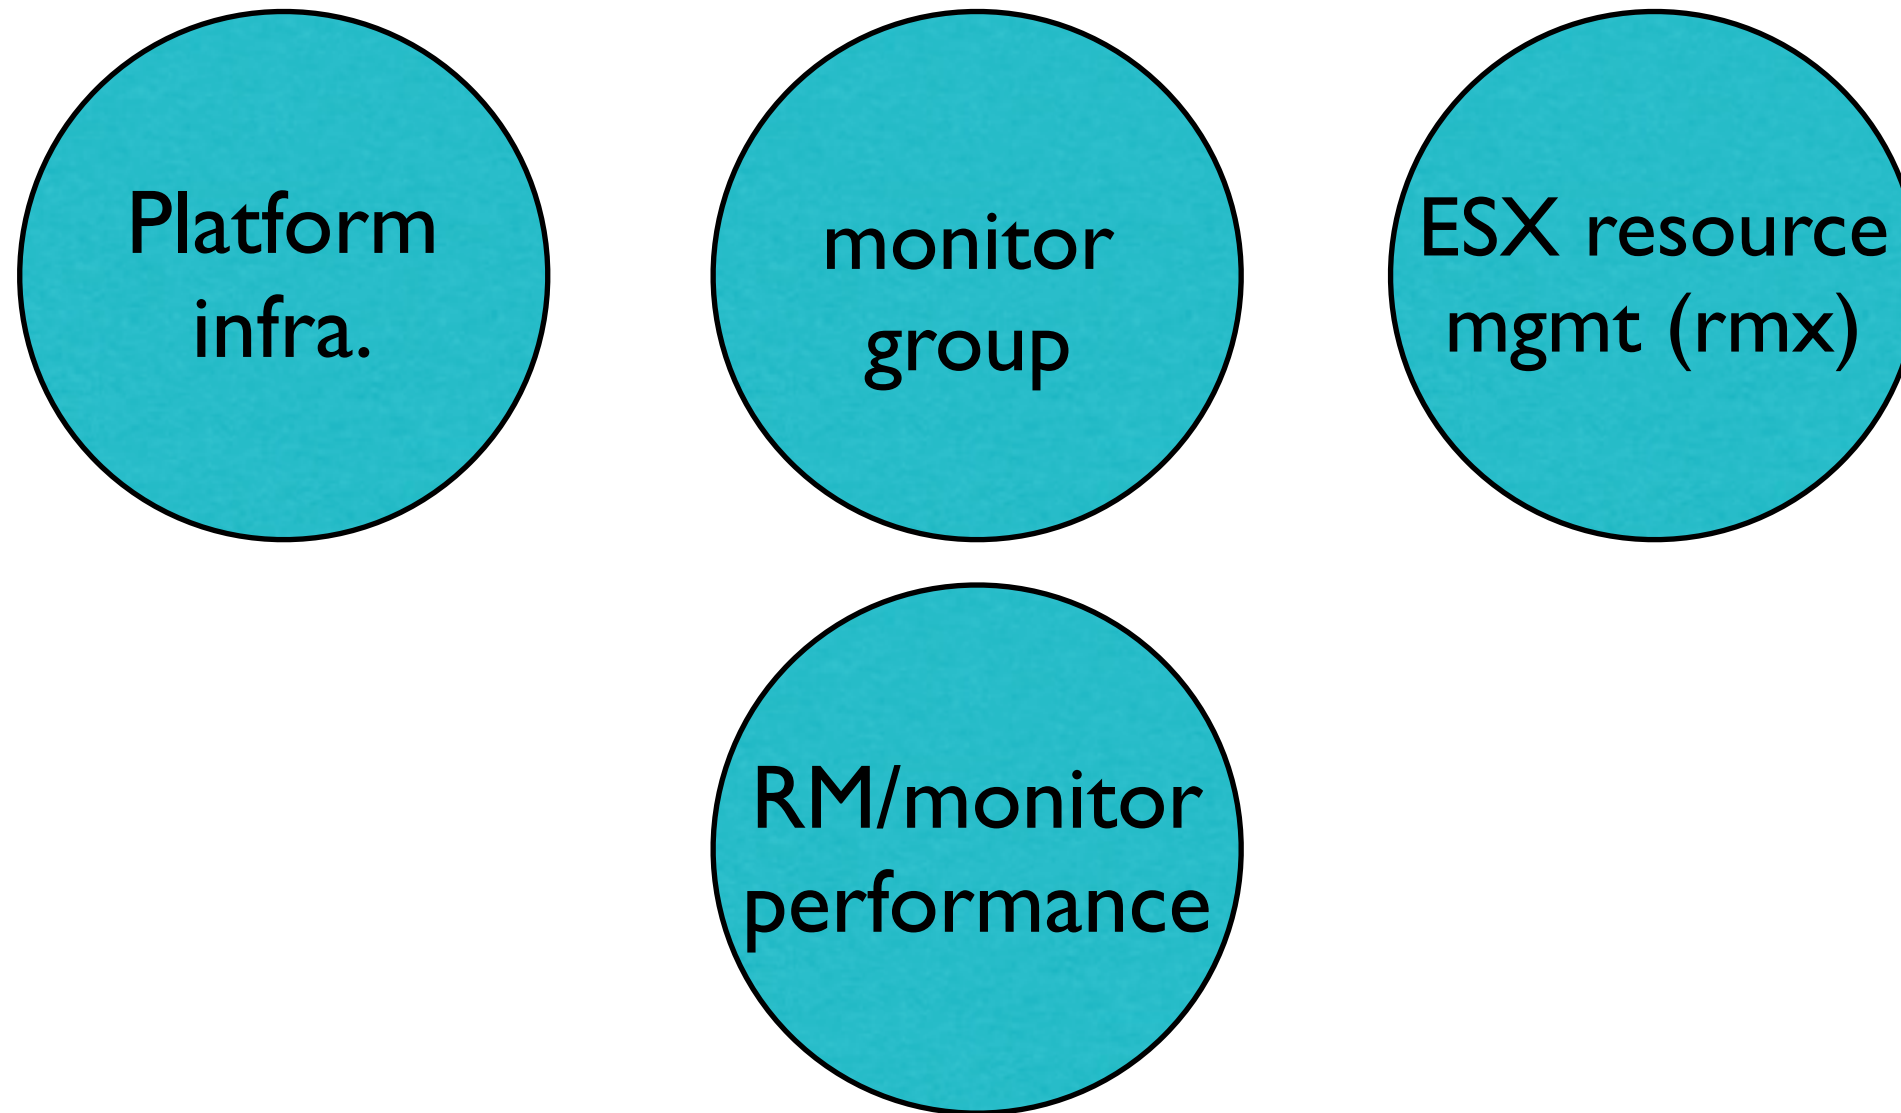
Low Memory State

- Correctness issue
- Monitor actions and repeated swapping
- Should be harder to enter low state
 - allocation and swap rates
 - avoid need to stop vcpus
 - work (in RMX): userworld swapping
 - not-exposed-to-vm

New Directions

- Distinguish translated from valid BPNs
 - improve locking/concurrency
 - eliminate stopping, avoid crosscalls
- Bulk-validation/invalidation
 - source of info. from guest and platform
- Restructure services
 - common structure/mechanism in vmm
 - tie reclamation/allocation rates
 - async page-sharing
 - selection, not-exposed pools
- New sampling techniques (miss-ratio curves)

Who's Involved?



vm-working-set@vmware.com

Mondays, 11am-12:30pm PDT, C13

Questions?


<https://wiki.eng.vmware.com/MonitorU>

monitor-list@vmware.com


Backups

Pre-KL Services


ballooning
(from guest)



sharing
(10Hz)




swapping
(swap action)




MemSched Services


ballooning
(from guest)



sharing
(10Hz)




swapping
(swap action)

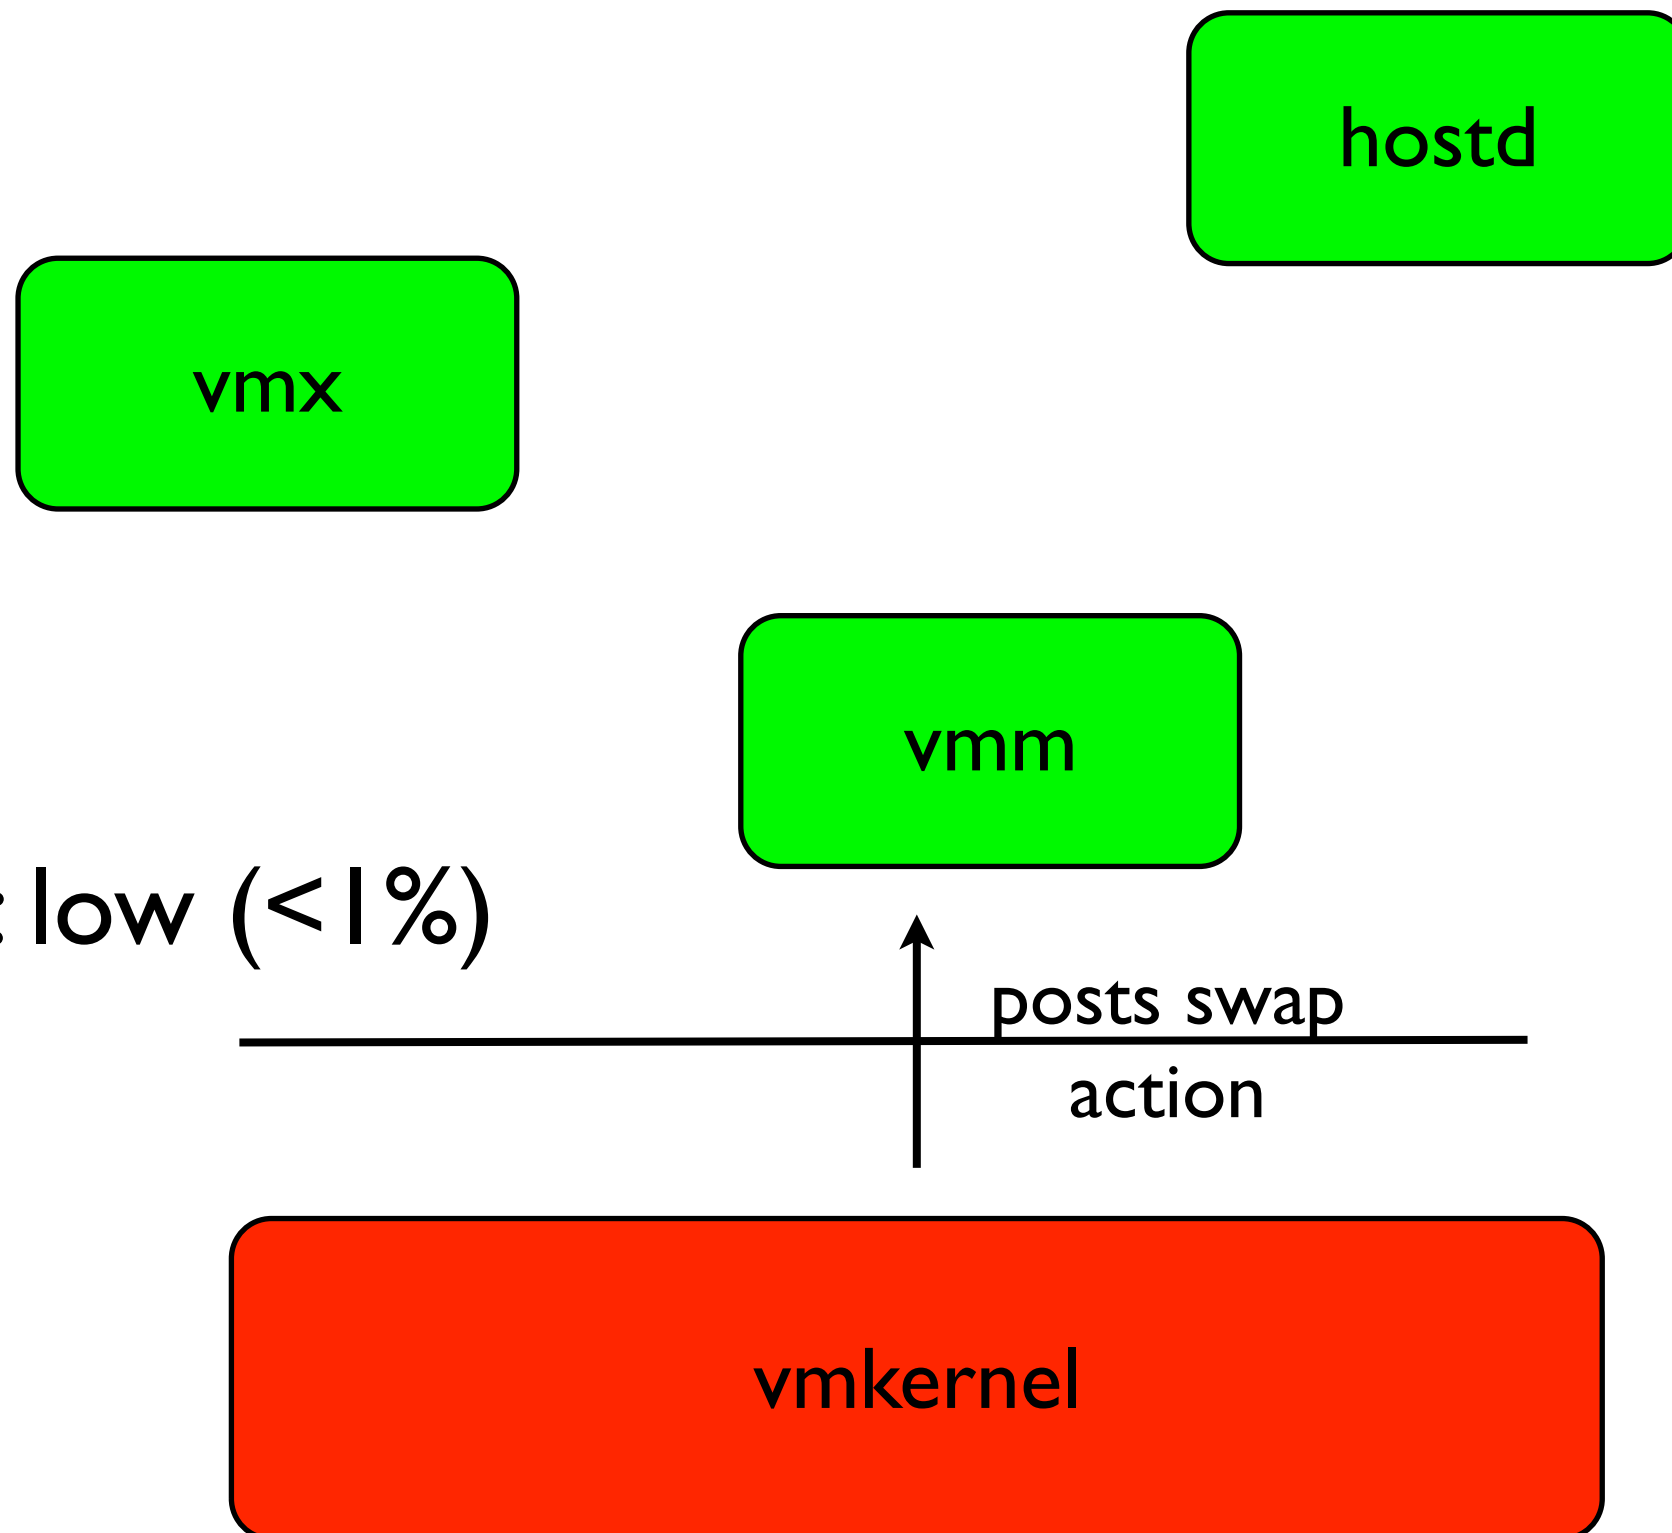


- knowledge of backing MPNs (swapping)
- cost of in-line computation (sharing)
- likelihood of access (sharing/swapping)

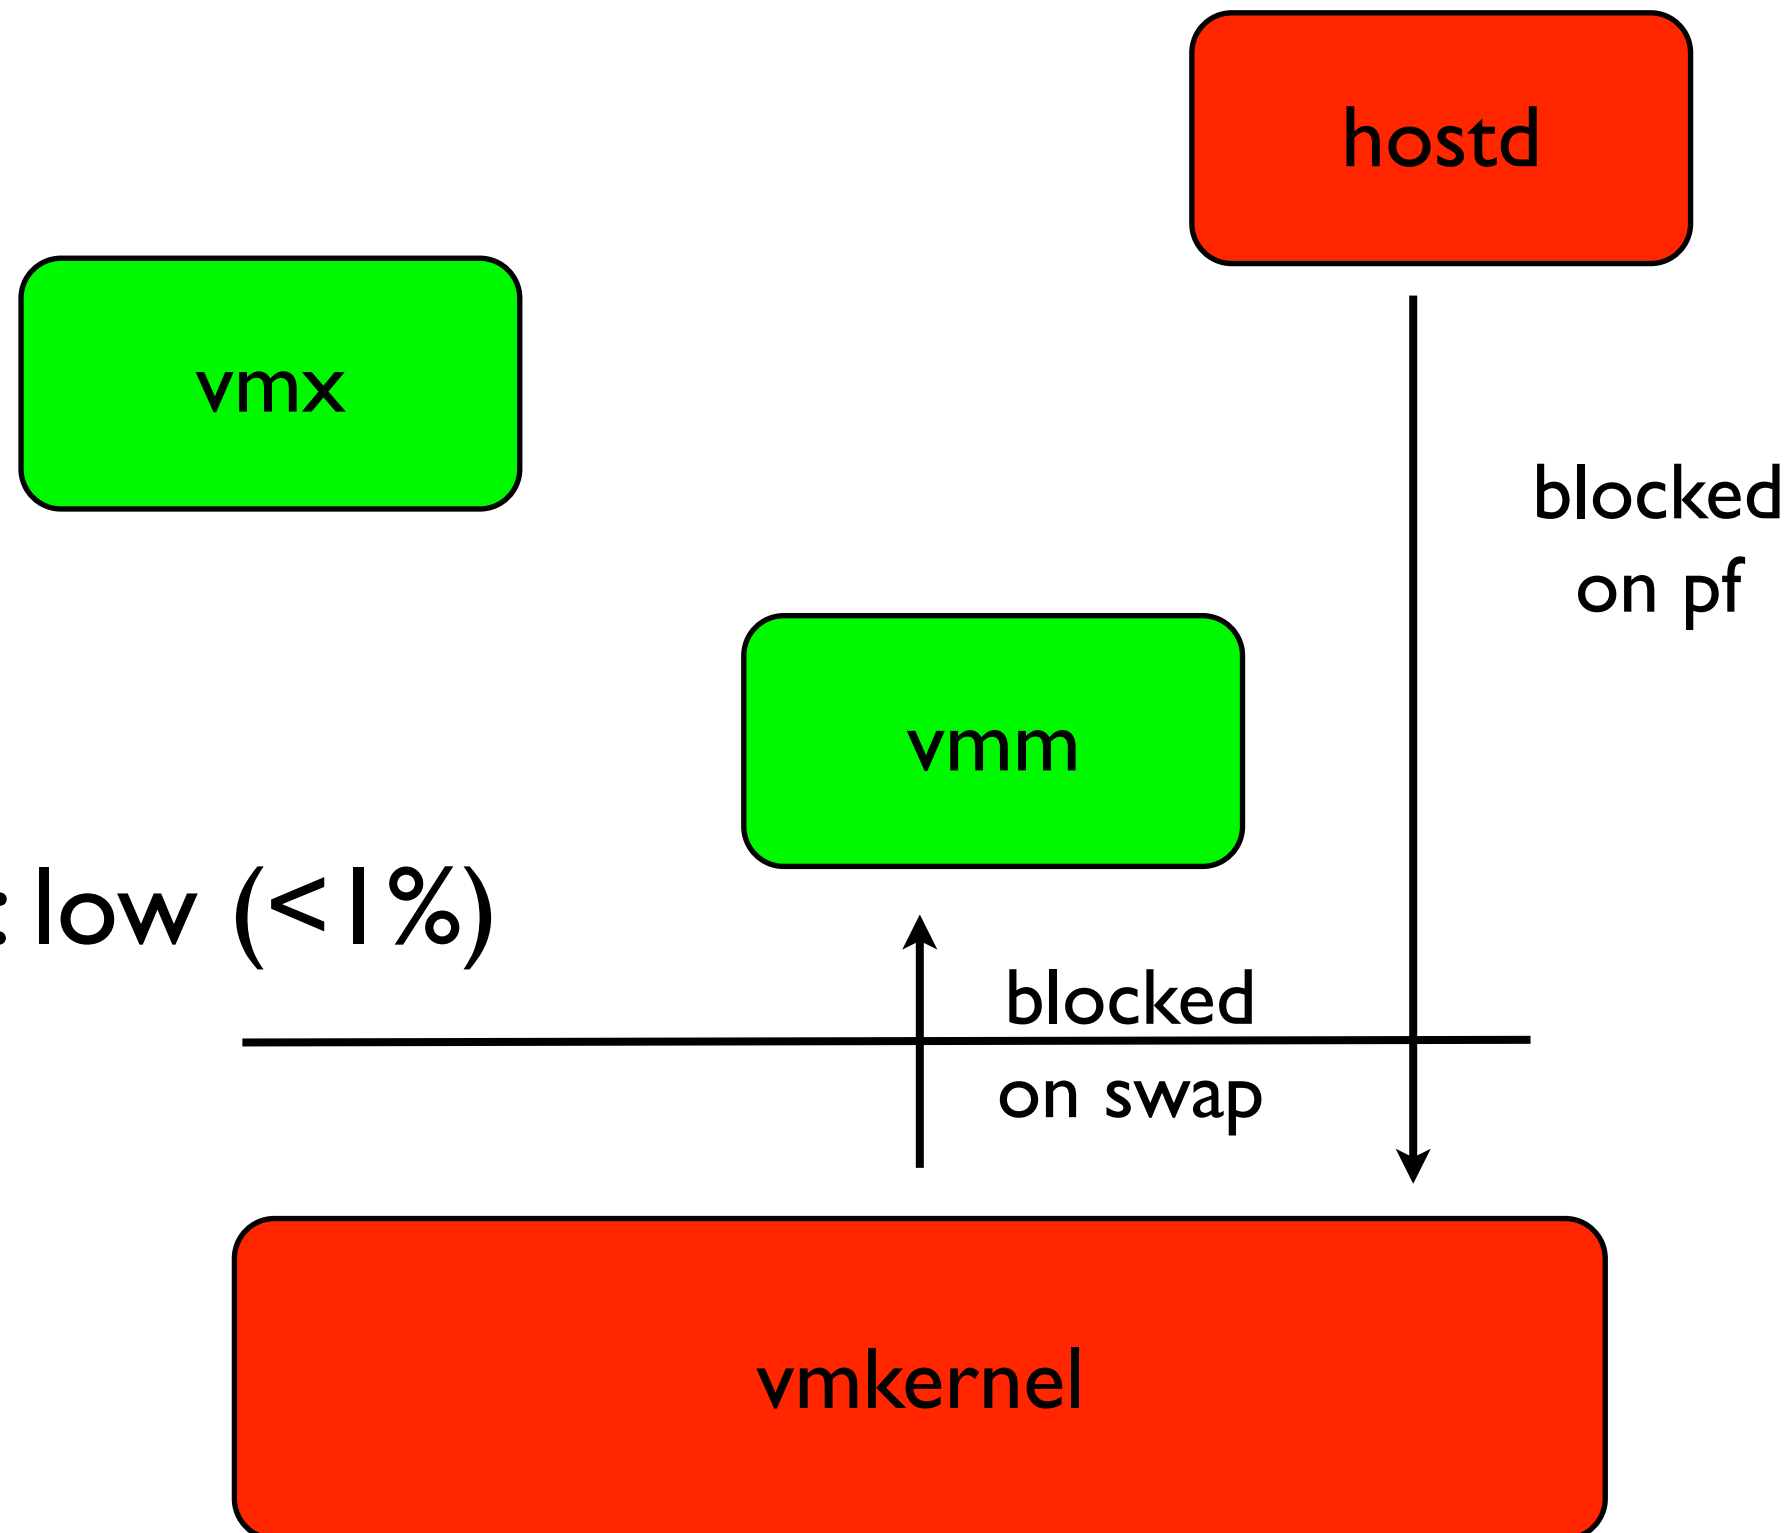
MemSched Deadlock



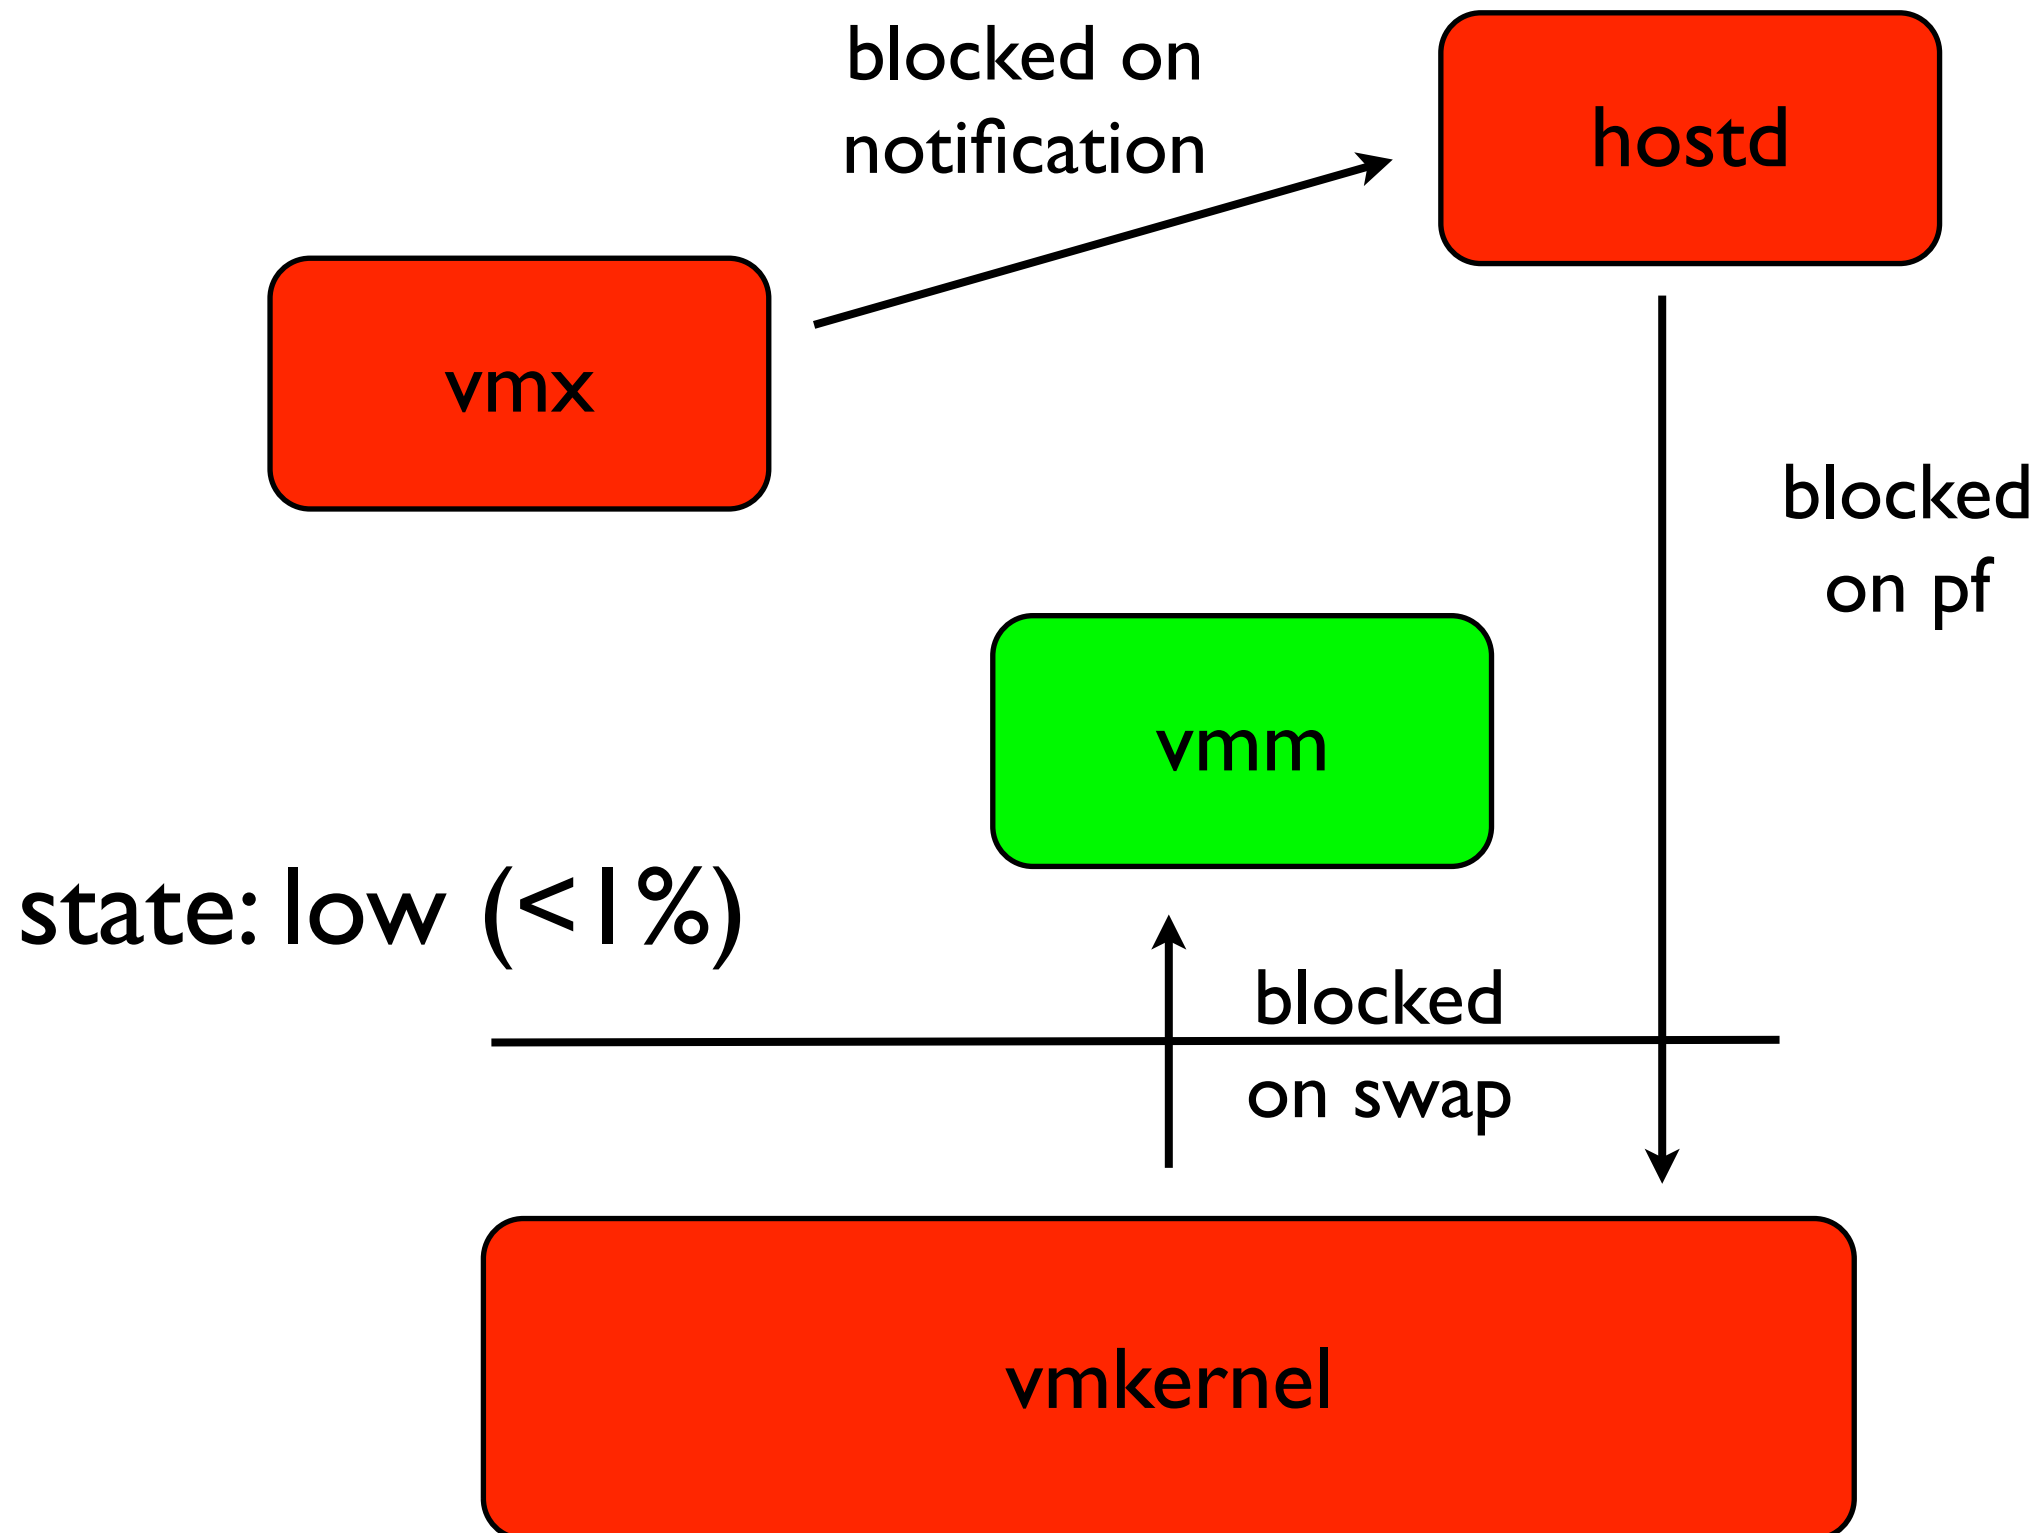
MemSched Deadlock




MemSched Deadlock



MemSched Deadlock



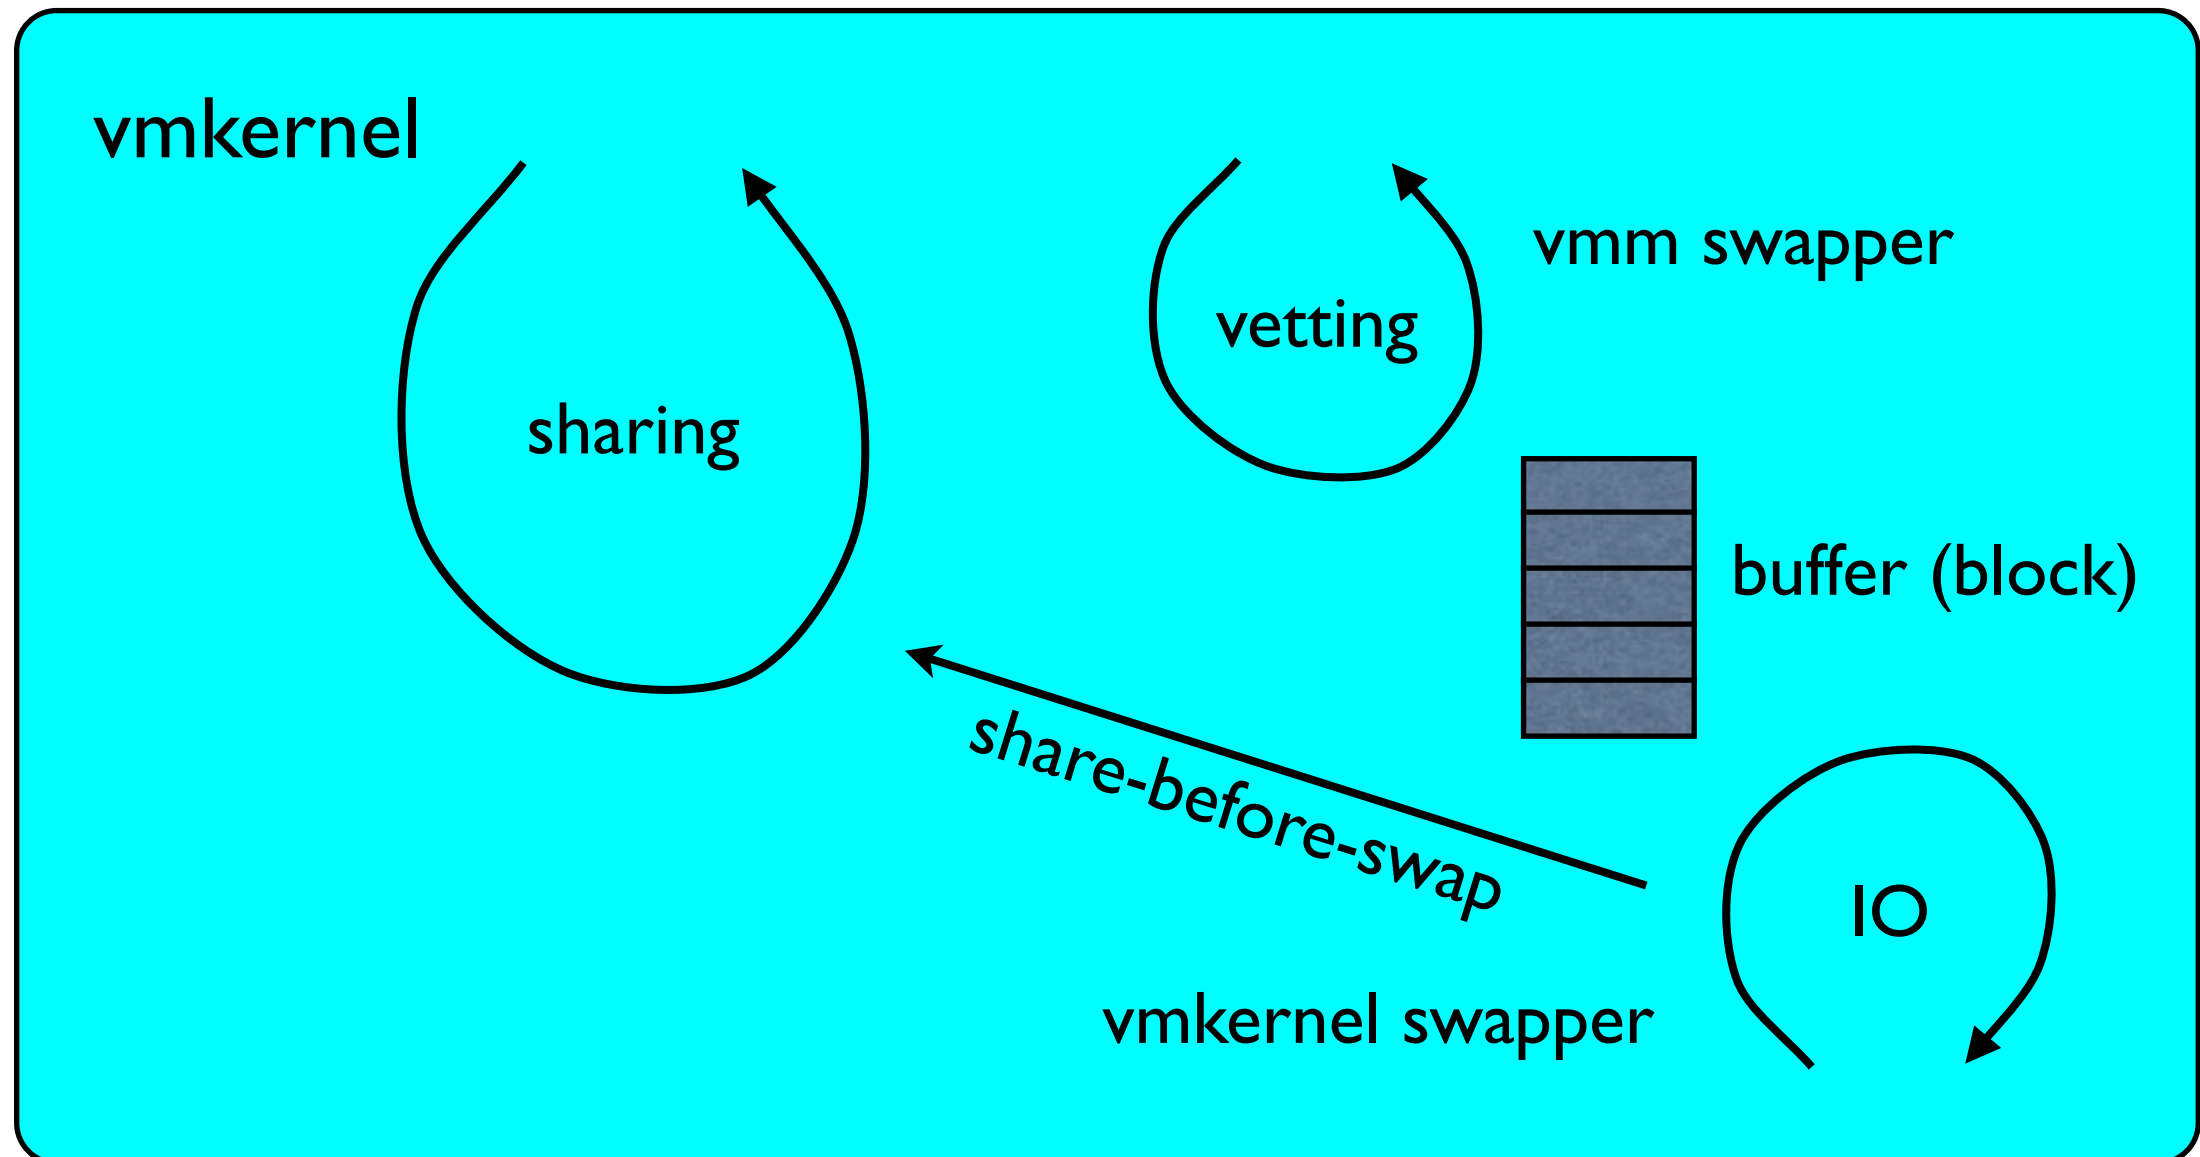
MemSched Deadlock



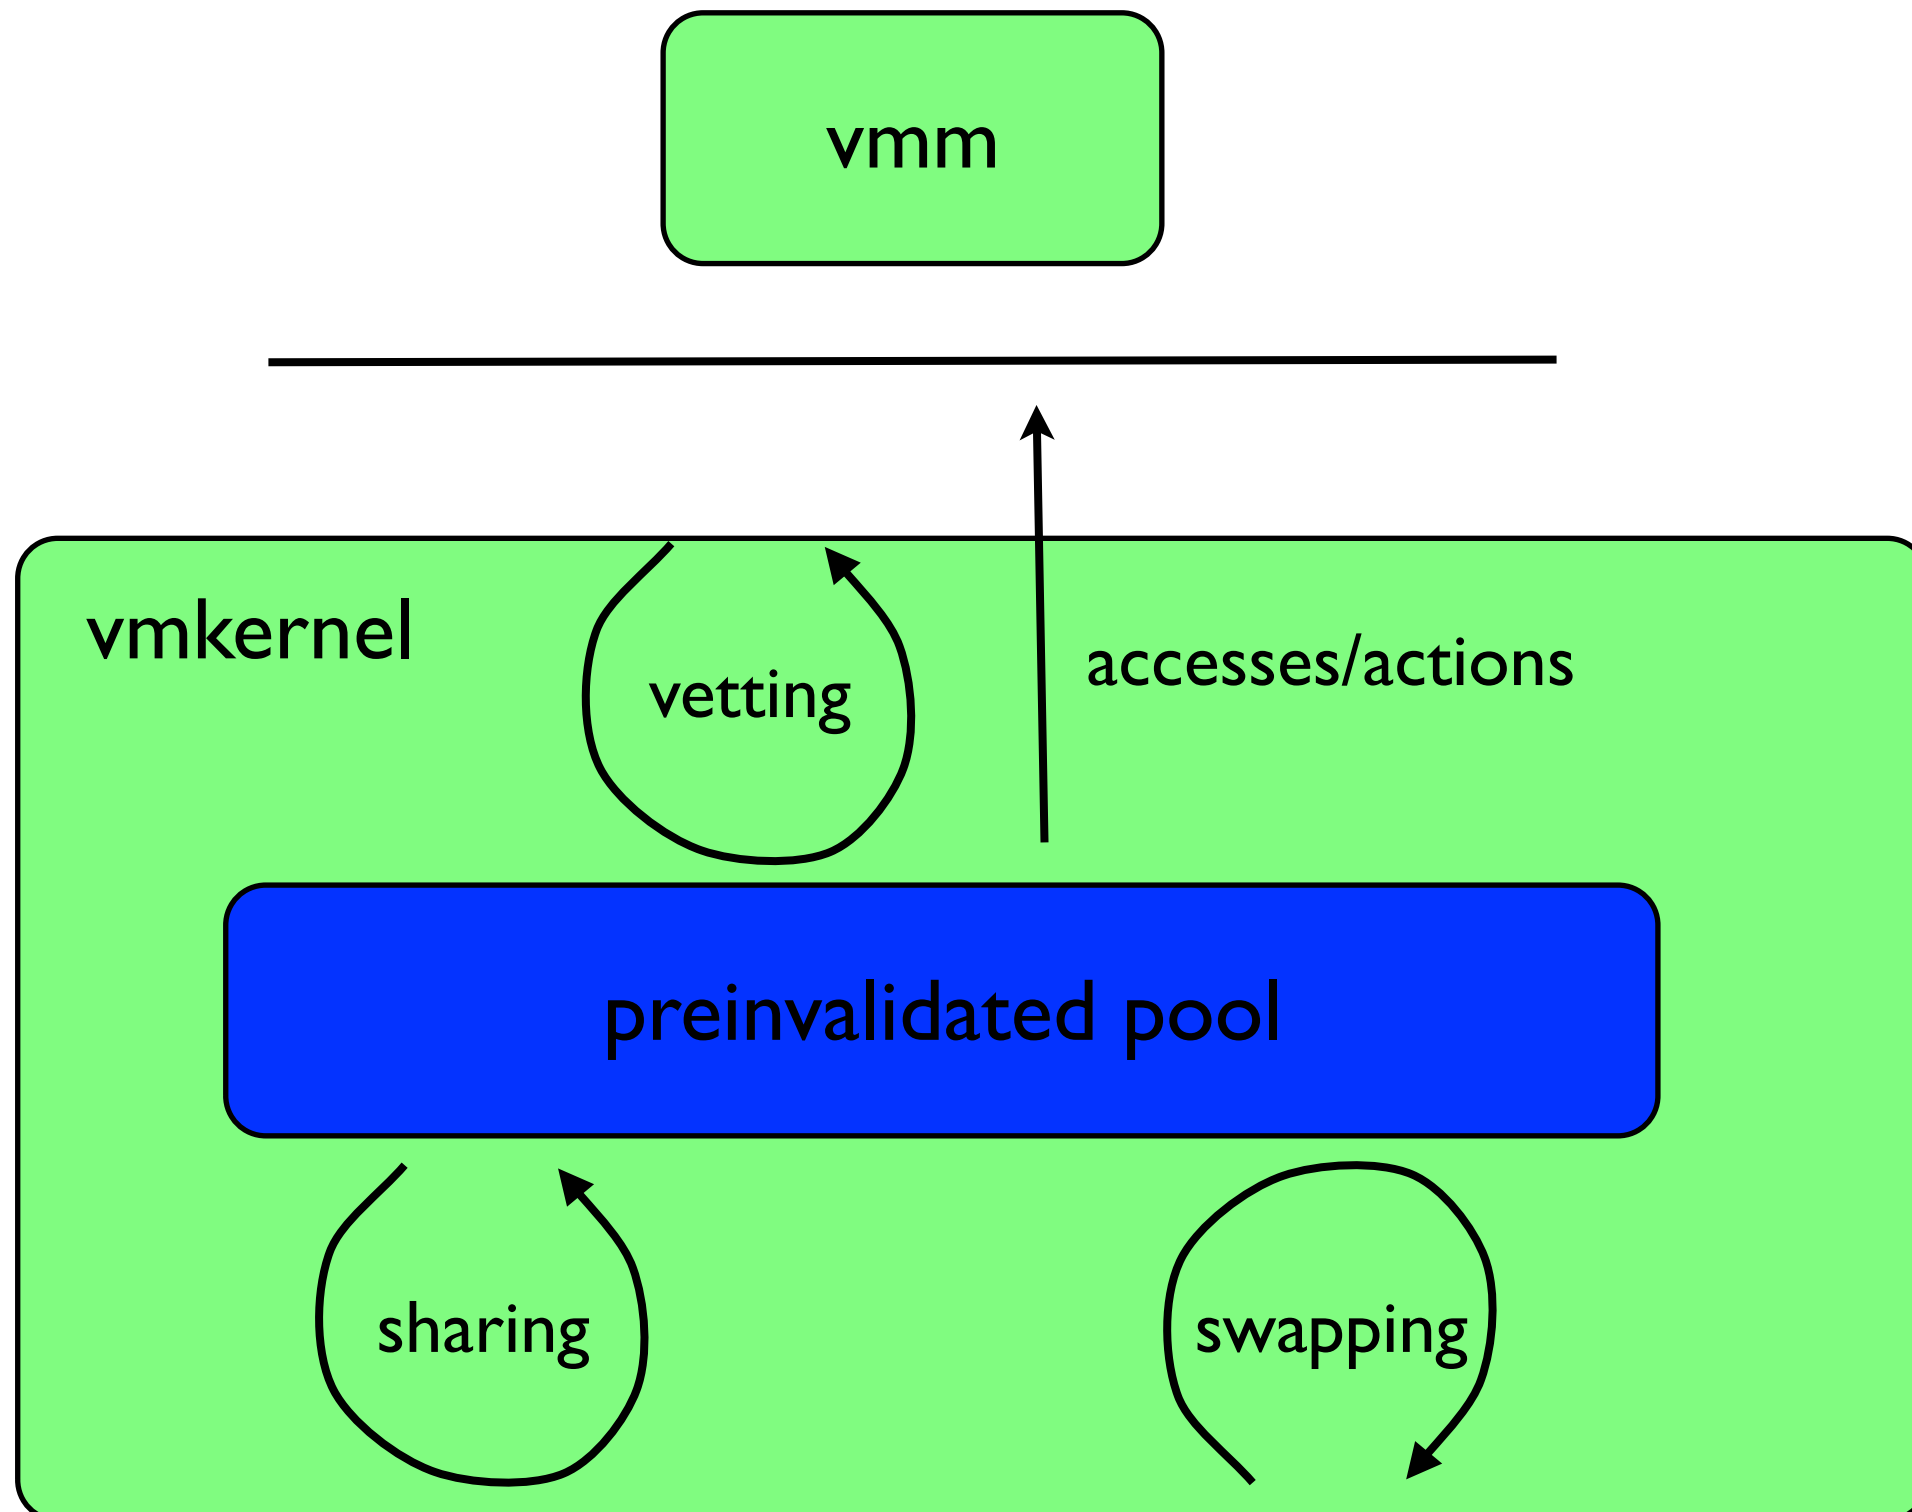
Many Solutions

- VMX-notification/vmkernel timeouts
- Allow invalidation across user-rpcs
- Vmm worker-threads
 - 1-vcpu VMs
- Privileging hostd worker threads
- Pre-invalidated/not-exposed pages

Services in vmkernel




MemSched Services



MemSched Services

vmm common action-driven structure



-
- extend to sampling, remapping (bulk validation)
 - other uses: simpler vmotion, trace-installaton

What's Different

- Separation of policy/mechanism
 - table-defined, common structure
- Removal of (expensive) synchronous work
 - limit on sharing rate
- Simplification of services within platform
 - vmm/vmk-swapper, sharing v. share-before-swap
- Pre-invalidation pool
- Remapping generalized to bulk-(in)validation
- Elimination of stopping to release memory
- Guest-hints to platform