

# Monitor and Guest Debugging

# Debugging

- - Monitor debugging
  - corequery for core debugging
  - debugstub for live debugging
- - Guest debugging
  - vmss2core for guest core debugging
    - leverages guests native kernel debugger
  - debugstub for live debugging
- - HW debugging
  - not covered

# Monitor core debugging example

- corequery provides very useful extensions
  - some must be run standalone
- Example: Backmapping anonymous MPNs to monitor VA's
  - Very useful for debugging corruption
  - Example: Find occurrences of 0xdeadbeef
- Could be automated - quicker but less fun

# Monitor core debugging example

```
- Use 'SearchForPat' to find PTEs with the target value:
# ./corequery open vmmcore-vcpu0 \; searchforpat 0xdeadbeef 0xffffffff \;
Found Pattern 0xdeadbeef in anon page 0x18f628 at offset 0xa8c
Found Pattern 0xdeadbeef in anon page 0x1d2887 at offset 0x864

- Dump monitor page tables (MON_PAGE_TABLE:0xffffffffc01c:0x20) and search:
(gdb) x/16384xg 0xffffffffc01c000

# grep 18f628 pt.txt
0xffffffffc01d920:      0x000000018f628061      0x000000011c7a8061

-Calculate mon VPN / mon VA for the PTE:
(gdb) p/x (0xffffffffc01d920 - 0xffffffffc01c000) / 8
$7 = 0x324
(gdb) p/x 0xffffffffc000000 + (0x324 << 0d12)
$8 = 0xffffffffc324000
(gdb) x/xw 0xffffffffc324000 + 0xa8c
0xffffffffc324a8c <FPU_PrepareX87PatchingForReplay+412>:      0xdeadbeef

-Repeat:
# grep 1d2887 pt.txt
0xffffffffc01dcb0:      0x00000001d2887061      0x00000001eddb1061

(gdb) p/x (0xffffffffc01dcb0 - 0xffffffffc01c000) / 8
$9 = 0x396
(gdb) p/x 0xffffffffc000000 + (0x396 << 0d12)
$10 = 0xffffffffc396000
(gdb) x/xw 0xffffffffc396000 + 0x864
0xffffffffc396864 <HVEExit+1300>:      0xdeadbeef
```

# Guest core debugging

- Use vmss2core to generate guest kernel dump
  - supports Linux, Windows, MacOS, Solaris, VMM/VMK, Elf, physical and more
  - finding guest symbols is half the battle
    - without them it's hard not to get lost
- Example Solaris:
  - On Linux run vmss2core:  
# vmss2core -S t-1244b748.vmss
  - generates unix.0 (kernel) and vmcore.0 (core)
  - In a Solaris VM run mdb:  
# mdb unix.0 (kernel) and vmcore.0

# Guest core debugging

- Dump the CPU info:

> ::cpuinfo -v

ID	ADDR	FLG	NRUN	BSPL	PRI	RNRN	KNRN	SWITCH	THREAD	PROC
0	ffffffffb304e0	1b	0	0	-1	no	no	t--8740328	ffffff0009a05c40	(idle)
		RUNNING	<--+							
		READY								
		EXISTS								
		ENABLE								

- Print the stacks:

> fffffff0009a05c40::findstack

stack pointer for thread fffffff0009a05c40: fffffff0009a05bc0

```
[ fffffff0009a05bc0 _resume_from_idle+0xf1() ]
  fffffff0009a05bf0 cpu_idle+0xbe()
  fffffff0009a05c20 idle+0x112()
  fffffff0009a05c30 thread_start+8()
```

- Print the log:

> ::msgbuf

MESSAGE

pcplusmp: pciexclass,060400 (pcieb) instance 3 irq 0x1b vector 0x31 ioapic 0xff intin 0xff is bound to  
cpu 0

PCI Express-device: pci15ad,7a0@18, pcieb3

- Use reference source code on OpenGrok to figure out what the guest is doing

# Guest core debugging

- Example Linux:
- On Linux run vmss2core:  
`# vmss2core -N6 VMwareRepro-d1a4aeac.vmss`
- generates 'crash' compliant vmss.core0
- Find linux kernel symbols (good luck)
- Find/build/run crash  
`# ./crash32 vmlinux-2.6.9-22.ELsmp vmss.core`

# Guest core debugging

- dump the log:

```
crash32> log
```

```
Linux version 2.6.9-22.ELsmp (bhcompile@porky.build.redhat.com) (gcc version 3.4.4 20050721 (Red Hat 3.4.4-2)) #1 SMP Mon Sep 19 18:32:14 EDT 2005
```

```
....
```

```
Unable to handle kernel paging request at virtual address bf3ca198
```

```
printing eip:
```

```
c0122406
```

```
*pde = a525dcbd
```

```
Oops: 0000 [#1]
```

```
SMP
```

```
Modules linked in: vmxnet3(U) vmxnet(U) ext3 jbd mptscsih mptbase sd_mod scsi_mod
```

```
CPU: 61
```

```
EIP: 0060:[<c0122406>] Not tainted VLI
```

```
EFLAGS: 00010086 (2.6.9-22.ELsmp)
```

```
EIP is at vprintk+0xe3/0x14a
```

```
eax: f7e0808b ebx: c040db8b ecx: 00020000 edx: 00007c94
```

```
esi: 00000246 edi: 0000000b ebp: 00000000 esp: f7e080f0
```

```
ds: 007b es: 007b ss: 0068
```

```
Process (pid: 0, threadinfo=f7e07000 task=f7e05000)
```

```
Stack: c1b55c00 00000007 c1b47c40 c0122320 f88c99a4 f7e08110 f88c140f f88c99a4
```

```
c1b47c40 c1b47c40 c1b55c00 f7e50d40 f88c169a c1b47c40 c1b47c40 c1b55c00
```

```
f7e50d40 c1b51b98 f88c617e c1b51b98 c1b55c00 f7e50d40 c1b51b98 c0222080
```



# Guest core debugging

```
- set the task
crash32> set 0xf7e31630
    PID: 1
COMMAND: "init"
    TASK: f7e31630  [THREAD_INFO: f7e08000]
    CPU: 11
    STATE: TASK_RUNNING (ACTIVE)

-dump the stack:
crash32> bt -S 0xf7e080f0
PID: 1      TASK: f7e31630  CPU: 11  COMMAND: "init"
#0 [f7e080f0] schedule at c02cf361
#1 [f7e0814c] __generic_unplug_device at c022207a
#2 [f7e08154] elv_requeue_request at c022086b
#3 [f7e0815c] scsi_request_fn at f88c61ac
#4 [f7e08170] __generic_unplug_device at c022207a
#5 [f7e08178] elv_requeue_request at c022086b
#6 [f7e08180] scsi_request_fn at f88c61ac
#7 [f7e08194] __generic_unplug_device at c022207a
#8 [f7e0819c] elv_requeue_request at c022086b
#9 [f7e081a4] scsi_request_fn at f88c61ac
#10 [f7e081b8] __generic_unplug_device at c022207a
#11 [f7e081c0] elv_requeue_request at c022086b
#12 [f7e081c8] scsi_request_fn at f88c61ac
...
#211 [f7e08f88] vfs_read at c0159c5e
#212 [f7e08fa4] sys_read at c0159e6f
#213 [f7e08fc0] system_call at c02d10c8
    EAX: 00000003  EBX: 00000003  ECX: bffff720  EDX: 000000e4
    DS:  007b     ESI: 00000000  ES:  007b     EDI: 00000000
    SS:  007b     ESP: bfff766c  EBP: bfff7828
    CS:  0073     EIP: 0804d03d  ERR: 00000003  EFLAGS: 00000292
```

# Live guest debugging

- Debug stub + GDB
- Add the following config settings:  
    `debugStub.allow=TRUE`  
    `monitor.debugOnStartGuest32 = "1"`

- Power on the VM

```
# vmware-vmx -qx VMwareRepro.vmx
```

```
...
```

Waiting for Guest 32-bit debugger to connect!

VMware Workstation is listening for debug connection on port 8832.

```
target remote localhost:8832
```

- Run GDB with the kernel symbols

```
# gdb vmlinux-2.6.9-22.ELsmp
```

```
(gdb) target remote localhost:8832
```

```
Remote debugging using localhost:8832
```

```
Debug stub remote connection accepted
```

```
[New thread 1]
```

```
0xffffffff0 in ?? ()
```

# Live guest debugging

- Set breakpoints:

```
(gdb) b panic
```

```
Breakpoint 1 at 0xc0121adf: file kernel/panic.c, line 60.
```

- Breakin and see where your guest is hung:

```
(gdb) bt
```

```
#0  get_8254_timer_count () at include/asm/io.h:396
#1  0xc039beb4 in wait_8254_wraparound () at arch/i386/kernel/apic.c:844
#2  0xc039bf22 in calibrate_APIC_clock () at arch/i386/kernel/apic.c:961
#3  0xc039bfe0 in setup_boot_APIC_clock () at arch/i386/kernel/apic.c:1000
#4  0xc039a91e in smp_boot_cpus (max_cpus=<value optimized out>) at arch/i386/kernel/smpboot.c:1085
#5  0xc0100487 in init (unused=<value optimized out>) at init/main.c:699
#6  0xc01041f1 in kernel_thread_helper () at arch/i386/kernel/process.c:256
```

- Allows debugging any guests from gdb on Linux

# Guest debugging with replay

- - very powerful tool to make guest failure reproduce deterministically
  - similar to PSMI and restart/replay for RTL debug at Intel
- - itrace and dtrace make guest debugging a mechanical process
- - limited to UP failures with no HMWWU
- Set a watchfault (optional) and start recording
  - `debug.watchfault.num=14`
  - `debug.watchfault.addr=0`

# Replay guest debugging

```
# vmrun beginrecording VMwareRepro.vmx rec1
StateLogger::STATE LOGGING ENABLED (interponly 0 interpbt 0)
StateLogger::LOG checksum
StateLogger::USING BOUNCE BUFFERS
Starting replayCheck
StateLogger::Continue sync while logging or replaying 27656
...
```

```
# vmrun endrecording VMwareRepro.vmx rec1
```

> Find the fault at address 0x160015 start itracing a few branches before this:

```
(5262472099) seq 23381963 DEBUG ( 2) len 224 brCnt 734656148 rip
0xc0105ca0 rcx 0xf7e07cd8 D___V___: DBG Fault# data 0xe000000:160015 lvl 2
kind 0
```

```
(gdb) x/i 0xc0105ca0
```

```
0xc0105ca0 <show_trace+17>: mov (%esi),%ebx
```

# Replay guest debugging

```
> Setup the following in your config settings and create a replay.cmd file  
replay.replayCmdFile=/root/replay.cmd
```

```
/root/replay.cmd:  
734650000 itrace 1  
734656200 itrace 0
```

```
> Start replaying:
```

```
# vmrun reverttosnapshot VMwareRepro.vmx rec5  
  
# vmware-vmx -qx VMwareRepro.vmx &  
[root@eng-rhel5-64 bug_666768]# PowerOn  
NUMA: automatic VM sizing request ignored  
VUsb powered on, but no USB controllers.  
Failed to load OpenGL utility (GLU) library.  
StateLogger::STATE REPLAYING ENABLED (interponly 1 interpbt 0)  
StateLogger::LOG checksum  
StateLogger::USING BOUNCE BUFFERS  
Restoring VMM  
SL: Resume (rip c0112696 holdoff 0)
```

# Replay guest debugging

```
> Examine itrace output and plan your next move - itrace or dtrace - to find out how XBX got 0x160015
StateLogger::
StateLogger::brCnt=734656145
StateLogger::eflags=0x2 CPL=0 codeSz=4 stkSz=4
StateLogger::XAX=0xc02dae86 XBX=0x160015 XCX=0xf7e07cd8 XDX=0xc02dae86 XDI=0x160015 XSI=0xf7e08000 XBP=0x68
XSP=0xf7e07cd8
StateLogger::R8=0x0 R9=0x0 R10=0x0 R11=0x0 R12=0x0 R13=0x0 R14=0x0 R15=0x0
StateLogger::ES=0x7b CS=0x60 SS=0x68 DS=0x7b FS=0x0 GS=0x0
StateLogger::cr0=0x8005003b cr2=0x6dba0ff0 cr3=0x37f8f2e0 cr4=0x6f0 cr8=0x0
StateLogger::
StateLogger::0x0060:0xc0105c96 | 89de                MOV    %esi,%ebx
MEMRD: LA0 c0105c98 Reg 1 PageNum 0x105 pgOff c98 size=1 CODE MEMRD: LA0 c0105c99 Reg 1 PageNum 0x105 pgOff c99
size=1 CODE
MEMRD: LA0 c0105c9a Reg 1 PageNum 0x105 pgOff c9a size=4 CODE
...
...
...
StateLogger::
StateLogger::brCnt=734656148
StateLogger::eflags=0x97 CPL=0 codeSz=4 stkSz=4
StateLogger::XAX=0x160ffd XBX=0x160015 XCX=0xf7e07cd8 XDX=0xc02dae86 XDI=0x160000 XSI=0x160015 XBP=0x68
XSP=0xf7e07cd8
StateLogger::R8=0x0 R9=0x0 R10=0x0 R11=0x0 R12=0x0 R13=0x0 R14=0x0 R15=0x0
StateLogger::ES=0x7b CS=0x60 SS=0x68 DS=0x7b FS=0x0 GS=0x0
StateLogger::cr0=0x8005003b cr2=0x6dba0ff0 cr3=0x37f8f2e0 cr4=0x6f0 cr8=0x0
StateLogger::
StateLogger::0x0060:0xc0105ca0 | 8b1e                MOV    %ebx, (%esi)

> Iterative process - challenge is doing it quickly

> itrace runs in the interpreter so monitor bugs may be converted to divergences.
```

# Debugging

- Tracking structures:
  - Page tracker/PMIO ring useful for debugging memory corruption
    - Hosted only \*
  - Ring buffer in the monitor useful for recent VM execution history and instrumentation
  - CCFs in BT mode provide record of recently executed guest code
- Questions/discussion ???