

Sampling Techniques

Alex Garthwaite

Monitor Group

December 19th, 2012

Confidential

vmware®

© 2010 VMware Inc. All rights reserved

Sampling in VSphere

- **Simple wss estimates of mainMem at BPN level**
 - access: memory rebalance
 - dirty: vmotion estimates
- **Reasons to sample**
 - memory reclamation: comparative measure of active
 - vmotion: estimate of dirtying rate/amount of work to do
 - VCOps, esxtop: measure of activity/efficiency of sizing of memory
 - DRS: anticipate expected memory need for placement decisions


Sampling in VSphere

- **Source code**
 - vmcore/vmm/main/busmemSample.c
 - vmcore/vmm/private/busmemSample.h
 - vmcore/vmx/main/busmemSample_user.c
- **Basic Framework**
 - Driven by 1Hz periodic callback
 - Based on vcpu progress (mean since vcpu0 skew)
 - Sample Period: 1 minute (ignoring first 4 seconds)
 - Sample History (number of sets): 4
 - Sample Size (per set): 100 pages
 - Each period:
 - update stats
 - kill old max-aged set
 - create new set of 100 pages
 - invalidate their mappings
 - VCPUs register touches when mainMem BPNs are translated to MPNs
- **Both vm configuration and VSI options to control settings**
 - needed? I think not

Sampling in VSphere

- **Start of Period**

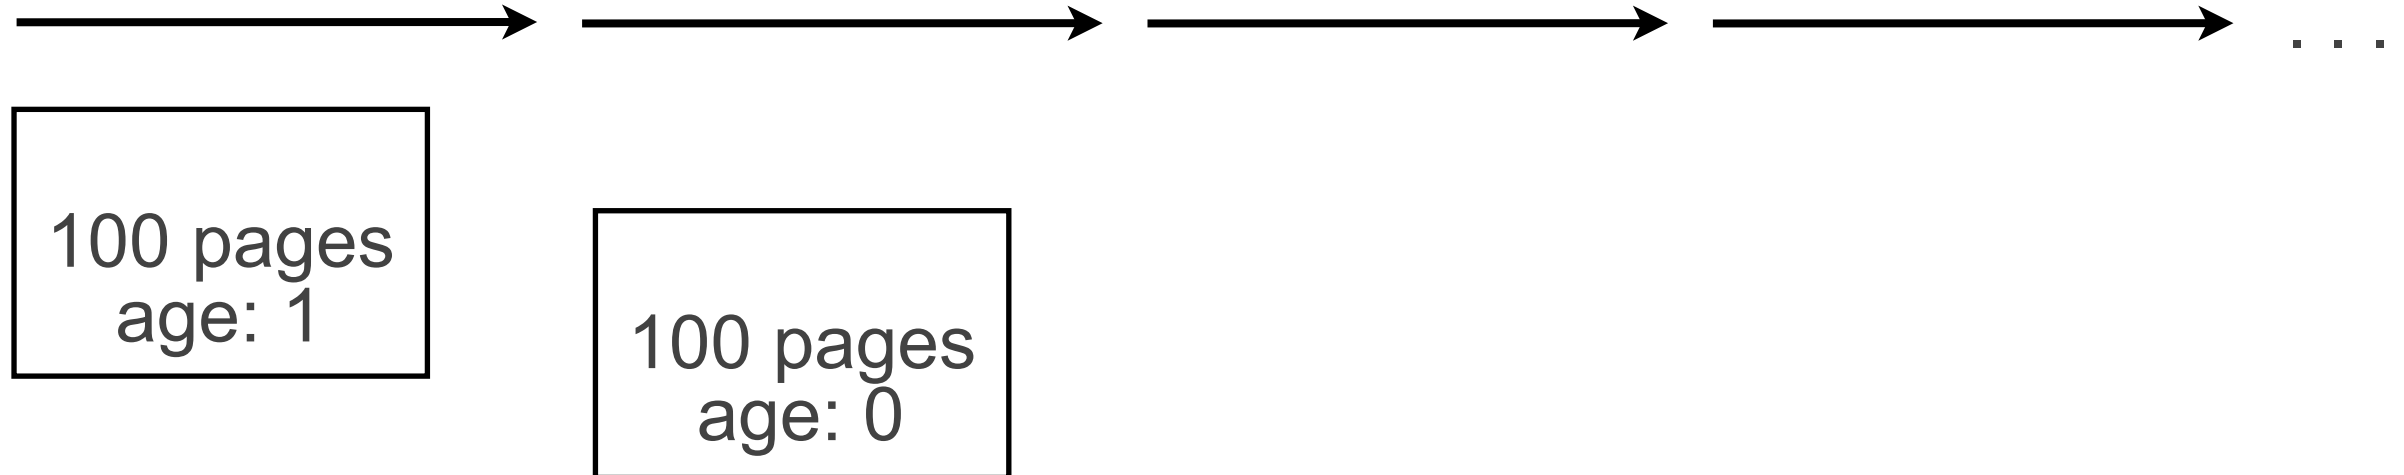
- concurrency
 - up to prod2013: stop all vcpus, use invalidation crosscall
 - now: use busmem lock and actions
- choose 100 pages
- invalidate mappings on all vcpus, even for pinned pages
 - first, global state: GPhys tables
 - in action, local state: MMU, etc.
 - leaves busmem frame/2MRegionInfo mapping state unchanged



Sampling in VSphere

- **Updating statistics**

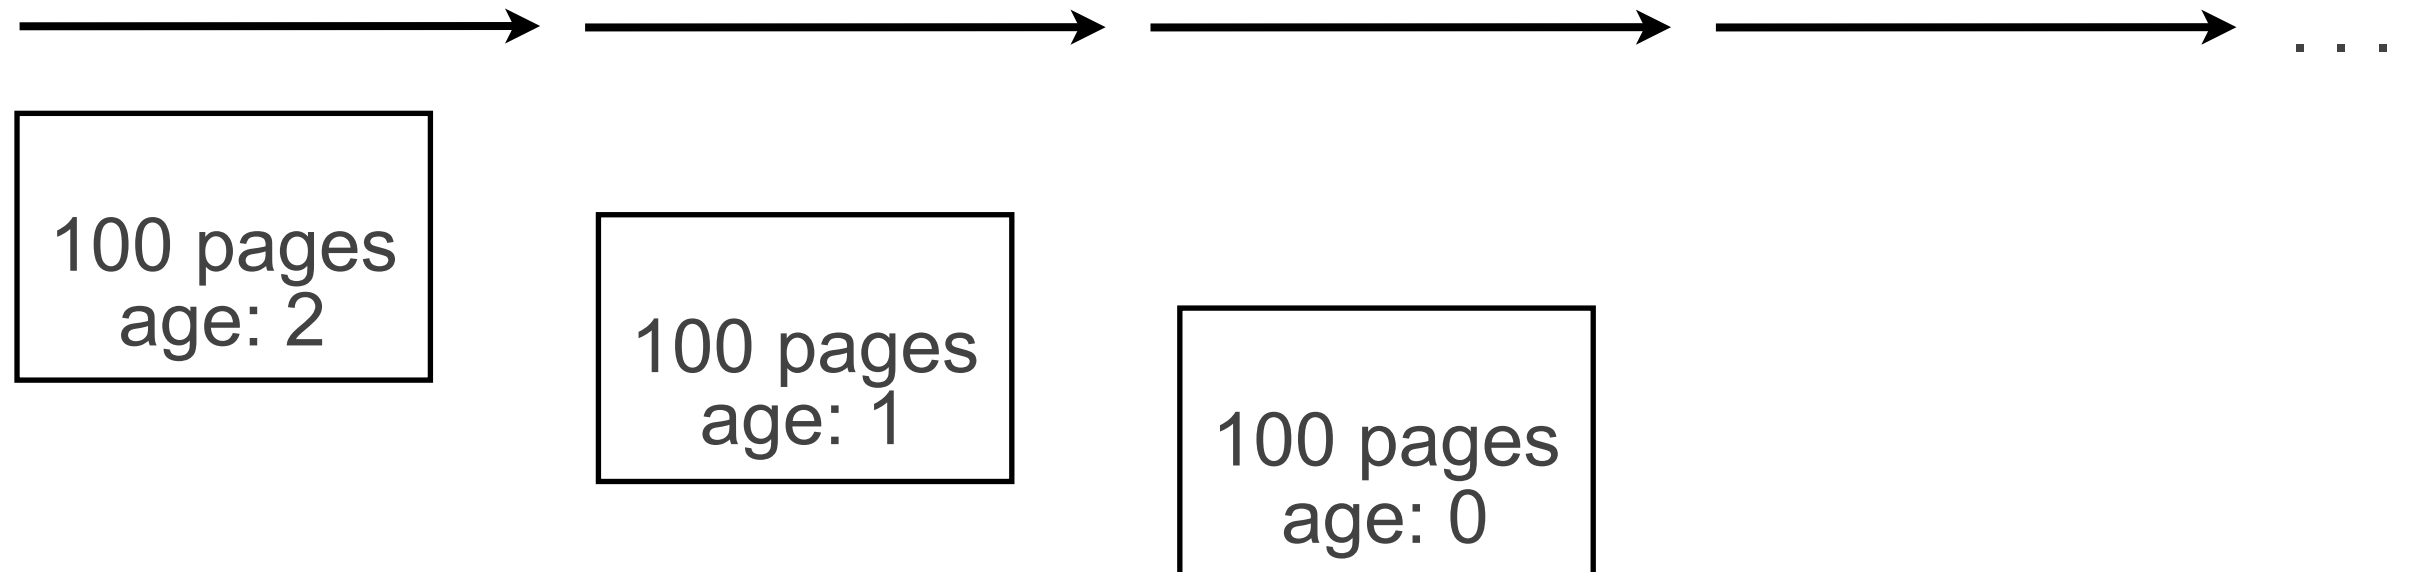
- for each active set, compute estimates of pct touched, pct dirtied
- for each age, aggregate fast and slow EWMAAs
 - fast weights: 75%, 25%
 - slow weights: 25%, 75%
- also: each (real) 1Hz Callback, measure fast/slow EWMAAs for newest set
- max of instantaneous, oldest, and youngest fast and slow averages used in memory balancing computations



Sampling in VSphere

- **Ballooning**

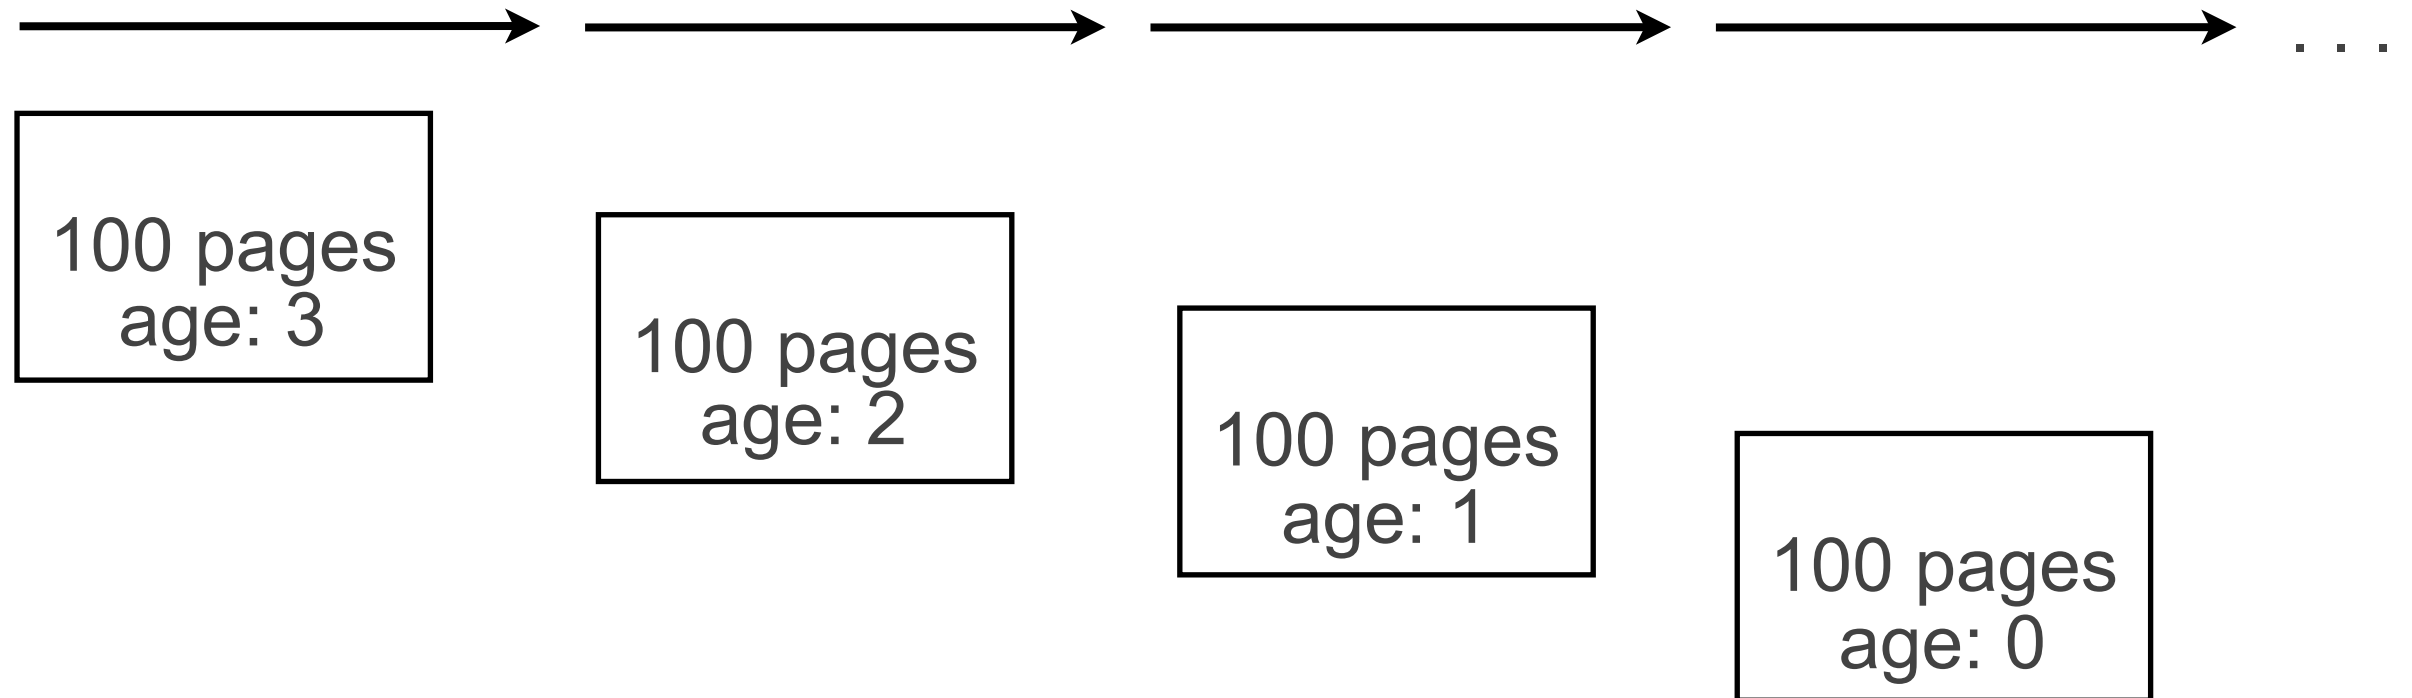
- may require guestOS to access/write sampled page
- one might preserve activeness values when not ballooned
- currently: clears touched/dirty sampled state when ballooned



Sampling in VSphere

- **Invalidation**

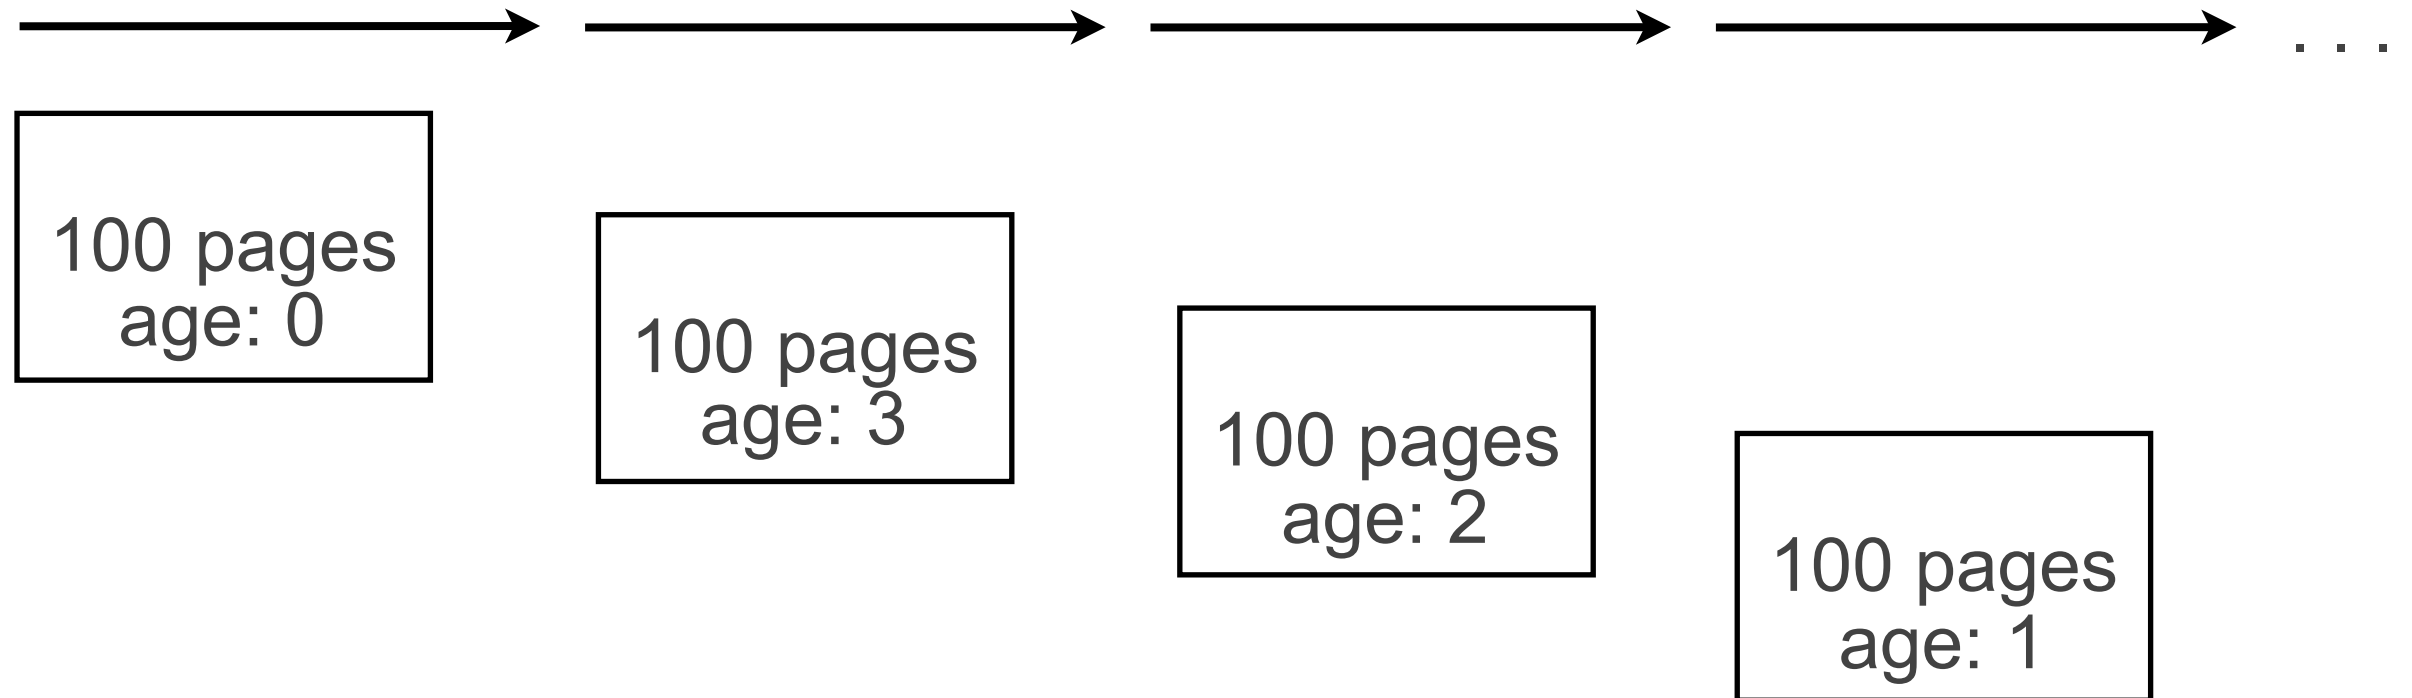
- under busmem lock, vcpu0 zaps GPhys for newly created set, posts actions
 - BusMemSampleInvalidateSampleSet, MONACTION_SAMPLE_SET, and BusMemSampleHandleInvalidations
- each set tracks which vcpus need to issue local flush
- in action handler under busmem lock, vcpus flush TLB and clear sw-mmuv state, etc.



Sampling in VSphere

- **Special Cases**

- (MMU) zero-pages considered touched/dirtied if selected
- invalidation must be done at rank 0 (MonTLB, pinned pages)



Common Questions

- **Why fault-based?**
 - A/D bits used elsewhere: would need mechanism to share uses
 - A/D bits preset as optimization (e.g., TLB prefetching, PTE update costs)
- **What do instant, fast, slow EWMA's mean?**
 - relative value
 - unclear how each maps to a time-window
- **Why 4 sets of 100 pages?**
 - relative error is inverse of square root of sample size, not population
 - impact on large pages
- **Why rotate sample sets?**
 - avoid poor (fixed) choice of pages skewing analysis
- **Should we support sampling for VMs with full reservations?**
 - alternative: use guest-level measures of activity in guest (assumes trust)

Large Pages and Sampling

- **Large mappings count as touching 4KB pages**
 - inflates wss estimates especially when used aggressively
 - sched.mem.alwaysTryLPageAlloc
- **Large mappings held off in presence of sampling**
 - preventing large mappings hurts performance
 - 2MRegionInfo structure tracks number of failed attempts
 - allow large mapping when
 - no sampled 4KB pages remain or
 - failed-threshold exceeded (512 attempts)
 - Note: failed-threshold only reset when 2MRegion translated large(!)
 - need to make rate: recorded/reset each Period
- **Sample set size**
 - worst case: every sampled page falls on a different 2M region
 - 400 pages can block 800MB from being mapped large

Large Pages and Sampling

- **Identifying large-page behavior**
 - separate A-bit walk of GPhys tables: identify breadth of access
 - wss-sampling identifies “cold” 2MB regions
 - currently: revalidates in GPhys as large mapping
 - need to evaluate if this is useful since 2MRegion wasn’t accessed much
 - can use failed-attempt count to categorize 2MB regions
 - consider multi-level sampling
 - sample N large pages, M small pages within each
 - use M (possibly as fixed pattern) to capture breadth

Per-VCPU Sampling

- **Alex Milouchev's VI-A project: location-aware sampling**
 - per-VCPU stats for page accesses
 - rearm sample-sets every 10Hz
 - able to determine which VCPUs access pages on which NUMA nodes
- **Overheads**
 - rearming pages: 2400x over lifetime of a set
 - 2-vcpu SPECjbb VM: increased translation by factor of 64x
 - overheads increased: ~8% or ~110 seconds over run
 - VMK_GET_PFRAME_MPN: +~93.7 seconds
 - VMK_LOCK_PAGE: +~15 seconds
 - BusMem_GetMPN
 - mapped bit still set for sampled pages but GPhys mappings gone
 - translation paths use BusMem_GetMPN
 - doubles #vmkcalls for sampled pages
 - Could fix BusMem_GetMPN to return INVALID_MPN for sampled pages
 - Could only call BusMem_GetMPN if page not destined for translation slow-path
 - Simplest is to just use BusMem_TryGetMPN

Latency-Sensitive VMs

- **Independently, Garrett found issues with the new barrier**
 - also saw high(er) BusMem_TryGetMPN costs
 - in BusMemInvalidateBackingStore
 - since most callers have zapped GPhys, unlikely to be useful
 - should just use INVALID_MPN in that code (especially for sampling)
 - saw serialization in the new action handler
 - Garrett's suggestion: post actions from vcpu to vcpu
 - alternative: don't require busmem lock in handler
 - only local state invalidated, similar to how zap-crosscall handler works
 - currently held to ensure set's BPN-list complete, for zero-page touching
 - could be done
 - should we consider not sampling highly latency-sensitive workloads
 - rely on guest metrics for right-sizing?

Miss-Ratio Curves and Marginal Benefits

- **Current metric too crude**
 - no notion of marginal benefit for adding or removing memory
 - idle-tax too coarse: leads to unfairness
- **Miss-Ratio Curves measure approx. fault-rate per allocation-level**
 - to contain cost, use combination of spatial and temporal sampling (RADIO paper)
 - track reuse-distances for pages in sample-set to build histogram
 - use histogram to compute number of accesses outside a given allocation level
- **Concerns**
 - overall cost of rearming every 10Hz
 - curve is flattened (especially toward the origin) through suppressing repeated accesses
 - access patterns within a period ordered by first access: slight lengthening of reuse-distances
 - sampling same # pages per VM or # pages proportional to memSize
 - how much each sampled page represents in comparisons
 - handling of resource pool hierarchy and aggregation: ignore for now
 - extend to, e.g., compression cache/swap to estimate larger allocations
 - unclear how to get feedback when ballooning
- **Should greatly improve memsched rebalancing**

Longer-Term Decisions: VCOps/DRS

- **Need longer term estimates of memory use**
 - current active estimate too crude/small window of time
 - needed for advising customers on right-sizing of VMs, placement
- **Rajesh's proposal**
 - consider possibly fixed set of pages
 - maintain reuse-distance histogram (64-bit counts)
 - each client can keep its own deltas for whatever time-scale it cares to track
 - need to checkpoint this histogram?

Current Status

- **Preeti moving sampling code to lib/sample**
 - plan is to move the sampling state to the platform
 - track in LockPage path when pages translated
 - can also capture vmkernel, vmx accesses
 - allows large sets of pages, also unifies sampling under memsched lock
- **Testing if in sample-set in vmm**
 - currently uses simple bloom filter to avoid updating any busmem state
 - Preeti's observation: clear dirty bits instead
 - busmem frames track dirtiness of 4KB pages: cache of platform dirtiness state
 - clearing this in vmm forces translations for write down slow-path (to LockPage)
 - would also need to switch to using BusMem_TryGetMPN to get MPN for non-write accesses
 - allows one to drop the bloom filter, simplify the translation slow-path test
- **PR957959 and sampling/GPhys interaction**
 - currently trying to verify if due to GPhys page-table walker taking "isLarge" flag

Recap

- **Fix sampling/GPhys bug (pr957959)**
- **Move sampling code to platform**
- **Implement/evaluate MRCs for memory balancing**
- **Improve sampling invalidation barrier**
 - eliminate use of busmem lock in handler
 - stop using BusMem_GetMPN in BusMem translation paths
 - stop using BusMem_TryGetMPN in BusMemInvalidateBackingStore
- **Use sampling activity to help categorize large-page use**
 - consider two-level selection at 2MB granularity
 - breadth, coldness measures
 - fix failed-attempt metric
 - consider dropping revalidation step for cold large pages
- **Consider not using sampling if latencies too high**