

VMMem Services

Alex Garthwaite

Monitor Group

December 18th, 2012

Confidential


vmware®

© 2010 VMware Inc. All rights reserved

Outline

- **PageList structure**
- **High-priority services**
- **Periodic/low-priority services**
- **Ballooning**

PageList Structure: new directions/64bit MPNs



- **Primary data structure for pages between vmm and platform**
 - status (success/fail VOID bit) encoded in 31st bit of BPNs
 - MPNs now only used for debug checks (vmm-platform matching) and remapping
- **Existing PRs**
 - eliminate MPNs from structure
 - stop using VOID bit and have BPNList and status
- **Other possibilities**
 - use single list at a time, not per-zap client
 - pass BPNs and statuses (or MPNs) separately
 - can stop populating MPNList during zap operation
- **Note: will need MPNList for pre/re-validation (e.g., remap)**
 - should evaluate benefits of ability to bulk-validate in vmm

New Framework for VmMem Services

- **Unify release/validation paths in the vmm**
- **Free platform scheduler**
 - move policy out of vmm
 - model: track and operate on not-exposed sets of pages (4K or 2M)
- **Address latency and overhead issues**
 - base work uniformly on guest progress
 - greatly reduce work in vcpu paths
 - track time stats to gauge latency bounds

Classes of Services

- High-priority
 - swap, p2mUpdate, page-retire
- Low-priority/best effort
 - NUMA-mig, 2M-defrag, page-sharing, ...
 - monitor/GPhys/frame mig, 2MB heuristics, page-checking
- Tension: hosted v. ESX (code & differences)

High Priority Cases

- Release operations (all invalidations):
 - swap: making progress wrt VMTrack
 - p2mUpdate: draining queue quickly/swap
 - page-retire: address failing HW
- Both polled and action-driven
 - responsive to allocation activity in vmm
- Do cases pending at start of check
 - vmotion, logging and action-starvation

VMM

```
while (PlatformGetRequest(&type, &numPages, &dataMPN)) {  
    services[type].processRequest(numPages, dataMPN);  
    PlatformCompleteRequest(numPages, dataMPN);  
}
```

VMkernel

- Tracks pending requests on first call, defers ones arriving after
- Selects type and page-set
- Operates on pages once not exposed to VMM

Implementation

- **Primary source files**

- VMM: vmcore/vmm/main/vmmem.c and vmcore/vmm/platform/*/platform.c
 - VMMemHandleReleaseRequests
- Vmkernel: vmkernel/mem/vmmemservices.c
 - VmMemServices_PickupRequest, VmMemServices_CompleteRequest
- Hosted: vmcore/vmx/main/vmmemServices.c
 - same as for ESX

- **Callbacks**

- VMM: vmmServicesCallbacks[type]
 - currently only one per type: filtering function
 - different across services:
 - p2mUpdate: only shared pages
 - swap: avoid pages currently in translation
- Platform: VMOMEM_SERVICES_CALLBACKS, VMOMEM_SERVICES_CALLBACK_FNS
 - selection, completion functions

- **Platform implicitly tracks first call for each type**

- list length and PageList in dataMPN passed back and forth

Where we are and clean-ups

- **Selection and dataMPNs**
 - a given service may require multiple iterations/batches to complete
 - swap selects all pages up front (limits max swappable target)
 - p2mUpdate selects pages from its own queue
 - should just select as the exchange of each batch happens
- **Pending tasks at start of handling release-request**
 - 32-bit atomic field: pending boolean indicating work
 - platform maintains own pending bitmap, copy passed back and forth
 - could merge these and simplify the code
 - makes tracking aggregate stats, timing info. simpler
 - does bake in order services are evaluated
- **Need for better (timing) stats than currently exist in vmm**

Low-Priority Services

VMM

```
IOHzCallback {  
    init(&timeBound);  
    while (PlatformGetRequest(&type, &numPages, &dataMPN, &timeBound)) {  
        services[type].processRequest(numPages, dataMPN, &timeBound);  
        PlatformCompleteRequest(numPages, dataMPN, &timeBound);  
    }  
}
```

VMkernel

- Prioritizes queue of requests based on timeBound
 - time since guest exit (ideal) or simple time-limit
 - E.g., sampling v. pshare v. remap-validation
- Selects type and page-set (avoid starvation)
- Operates on pages once not exposed to VMM
 - On-demand accesses by guest handled per page

Low-Priority Cases

- Invalidation
 - NUMA-migration, 2M defragmentation
 - page-sharing
 - 2M-region promotion
- Validation (needed? need perf evaluation)
 - NUMA-migration
 - checkpoint restoration
 - based on guest access patterns
- Sampling: possibly leave out

Low-Priority Cases

- Tracking why pages are not exposed
 - handle in page-fault path
 - alternative proposal: mix of in-line and async work
- Queues of work (and prioritization)
 - monitor-side operations
 - platform-side operations

Timing

- **Latency requirements**
 - `memschedInfo.vmmem.prop.latencySensitivity`: 64-bit value!
 - need map to timing constraints
- **Overall time**
 - time since guest exit too expensive
 - measure time since vcpu-progress stops as approx.
- **Current 10Hz period**
- **Per-type timing stats**
 - easiest to measure from vmm
 - useful in platform to choose work
 - ... but may not be issue: little in-line work

Prioritization

- **Monitor and Platform managed queues**

- split time bound between two, interleave

- **Starvation concerns**

- may not be an issue given new structure of services: work out-of-line
- groups of types' pages can be batched
- timing overall may be short: crosscalls/platform-calls dominate

- **Strawman proposal**

- sampling > {migration, defrag, sharing} > {validations} > 2M-promotion

Initialization

- Configuration
 - latency limits
 - timebound for this round of processing
- Timing info
 - time exited guest
 - time this round started

Request

- Request state (op, etc.)
- Stats
 - time stats per operation type

Asynchronous Activity In Platform

- **Pages passed down not exposed**
 - need to allow page-faults before processing complete
 - may cancel operation on that page: swap-outs
 - may need to complete for those pages: page-sharing, remapping
 - requires knowing *why* page is being processed (easy)
- **Extensions**
 - with less impact on guest, re-evaluate page-sharing, NUMA-mig rates
 - base more on memory bandwidth
 - remapping at 2MB level
- **Accounting/scheduling asynchronous work**
 - basing on VM's allotment may result in starvation of service
- **(Re)validation services may add overhead through buffering**

Ballooning

- **Measuring guest impact**
 - long-standing problem: measuring impact and backing off
 - dynamically changing boundary: not a static decision, how to handle this?
 - use guestLib: access to pinned info., etc. often easier from user-level in guest
 - learn from vmmon/hosted environment and improve *both*
- **Balloon targets and role of ballooning**
 - primarily constrains guest physical address space
 - secondarily releases memory (after guest execution to allocate pages)
 - so target may be much higher than perceived level of overcommit
 - allows some applications/guestOSs to constrain cache structures, etc.
 - what of guestOSs that maintain free-page pools (e.g., 10% of memsize in Windows)
- **Downward spiral (in general)**
 - previously: sampling bias (only tracked non-ballooned pages)
 - maintained wss percentages then applied to current non-ballooned size
 - guest less active, not on normal tasks when ballooning or swapping
 - impacts both pages touched and vcpu progress
 - what if free-page pools, etc. not all scaled down?
- **Fairness issue across VMs**
- **Will never solve problem of VMs too early in power-on cycle**

Recap

- **Userworlds/VMX (from yesterday)**
 - numa-migration support
 - backing with 2MB pages
 - better page-sharing?
- **High-priority services**
 - simplify both data (selection) and control (pending) aspects
 - add timing/stats
 - add page-retirement zaps (if not there yet)
 - bound VMX pinning (p2mUpdates)

Recap

- **Low-priority services**

- redirect appBallooning queue to platform
- services
 - split remap into two halves: invalidation, validation
 - evaluate need for (remap) validation (TPC-C runs on NUMA)
 - move page-sharing to new framework
 - consider moving 2M-defrag to this framework
 - consider prevalidation in vmm, “rearming” large regions
- implement asynchronous platform code
 - support 2MB
 - dealing with page-faults on in-process pages
- evaluate timing of services and how throttled
- introduce (latency-aware) timing model
- add better stats/timing info.
- in vmm and platform, balance across pending queues to avoid starvation