

VProbes Update

What's New Since ESX 5.1

Outline

Language Features

- Exceptions
- Loops, iterators
- Binary logging
- Aggregate formatting
- Shared integers
- Preprocessing
- Probe globbing

Fault Injection

- Modify register state
- Modify memory
- Skip functions
- Inject virtual MCEs

Usability Enhancements

- External types
- Local variables and arguments
- Line number probing

Language Features

Exception Handling

■ Intercept runtime errors

- Typically builtin failures
- Memory read errors, full bag on insert, empty bag on remove, etc.

■ Programatically handle errors

- Trap and ignore errors
- Print error message
- Assign default values

■ Syntax similar to C++

■ Default behavior:

- Terminate the current probe
- Print an error message

```
try {  
    value = *address;  
    printf("%d", value);  
} catch (VmkMemAccessError) {  
    printf("Memory read error\n");  
}
```

Loops

- Previously:
 - Iterative computations expressed using recursion
- Standard loop statements now available:
 - For, while, do-while
 - Break and continue

- Example:

```
memmodel guest64;  
void guestbt() {  
    printf("%#x\n", RIP);  
    for (rbp = RBP; rbp; rbp = *(uint64*)rbp) {  
        printf("%#x\n", *(uint64*)(rbp + 8));  
    }  
}
```

Bag Iterators

- Iterate over all key-value pairs in bags
- Syntax inspired by Python
- Example:

```
void printbag(bag b) {  
    int key, val;  
    foreach (key, val) in b {  
        printf("%d: %d\n", key, val);  
    }  
}
```

Avoid Infinite Loops

- Bound time per probe fire
- Older releases:
 - Upper bound on function calls per probe (400)
- Newer releases use a cost model:
 - Assign costs to each builtin, loop back-edges, function calls
 - Impose a maximum budget per probe (aimed at about 50us per probe)

```
try {  
    for (i = 0; ; i += 1) {}  
} catch {  
    printf("Loop ran for %d iterations\n", i);  
}
```

Multiple Returns

- Multiple return points in functions and probes
- Return different values:

```
int squareSum(int n) {  
    if (n == 0) {  
        return 0;  
    } else {  
        return n * n + squareSum(n - 1);  
    }  
}
```

- Early probe returns

```
VMK1Hz {  
    if (PCPU == 0) {  
        return;  
    }  
    ...  
}
```


Binary Logging

- `printf` outputs formatted text
- `writeint` outputs binary data

`writeint` (uint64 data, int numBytes)

- Provides some form of output compression
- Use repeated calls to write out entire structures
- Possible binary output formats:
 - Fixed-size output records
 - Variable-size records (must also write record length)
- Problem: difficult to distinguish errors from binary data
 - Vprobe output split into output and error streams
 - Stdout/stderr on ESX; vprobe.out/vprobe.err on hosted

Formatting Aggregate Output

■ Aggregates

- Data type for building Histograms, bucketed by integer and string keys.
- Print using `logaggr(a)`, clear using `clearaggr(a)`
- Example:

```
aggr ioStats[1][1];
bag ioSize, ioStart, ioType;
FS3AsyncIOStart(int len, void *o, int read, int guid) {
    ioSize[guid] <- len / 1024;
    ioStart[guid] <- TSC;
    ioType[guid] <- read;
}
FSAsyncIOEnd(int guid, int status) {
    ioStats[ioSize[guid], ioType[guid] ? "READ" : "WRITE"]
        <- cycles2ms(TSC - ioStart[guid]) / 1000;
}
VMK1Hz logaggr(ioStats);
```

Formatting Aggregate Output - logaggr

- `logaggr(ioStats)`
 - Fixed format, tabular.
 - Each row is a bucket in the aggregate.

| intKey0 | strKey0 | avg | count | min | max | pct% |
|---------|------------|-----|-------|-----|-----|-------|
| | 0x2c READ | 7 | 92 | 0 | 210 | 0.4% |
| | 0x18 READ | 23 | 31 | 0 | 273 | 0.4% |
| | 0x40 WRITE | 45 | 18 | 0 | 208 | 0.5% |
| | 0xc WRITE | 38 | 32 | 0 | 89 | 0.8% |
| | 0x8 WRITE | 32 | 38 | 0 | 87 | 0.8% |
| | 0x4 WRITE | 7 | 165 | 0 | 158 | 0.8% |
| | 0x4 READ | 7 | 170 | 0 | 292 | 0.8% |
| | 0x60 WRITE | 30 | 46 | 0 | 210 | 0.9% |
| | 0x8 READ | 23 | 69 | 0 | 175 | 1.1% |
| | 0xc READ | 25 | 65 | 0 | 218 | 1.1% |
| | 0x80 READ | 5 | 2482 | 0 | 727 | 8.8% |
| | 0xa0 WRITE | 50 | 2127 | 0 | 308 | 72.0% |

Formatting Aggregate Output – `printa`

- `printa(aggr[, format-str])` [Available in 2014]
 - C `printf` style format string “applied” to each bucket in the Histogram before printing.
 - Arguments to the format string are implied,

"%x %x .. %s %s %d %d %d %d %d"

(intKey0 intKey1 .. strKeyN-1 strKeyN avg count min max pct)

- **Explicit fields selectors:** “%<field-identifier>\$<format-specifier>”
 - Inspired by Posix extensions to `printf` for numbered arguments.
 - **Field Identifiers:**
 - %0\$, %1\$, %2\$...
 - %avg\$, %count\$, %min\$, %max\$, %pct\$

Formatting Aggregate Output - Examples

■ `printa(ioStat, "%d")`

| intKey0 | strKey0 | avg | count | min | max | pct% |
|---------|---------|-----|-------|-----|-----|-------|
| 96 | WRITE | 30 | 46 | 0 | 210 | 0.9% |
| 8 | READ | 23 | 69 | 0 | 175 | 1.1% |
| 12 | READ | 25 | 65 | 0 | 218 | 1.1% |
| 128 | READ | 5 | 2482 | 0 | 727 | 8.8% |
| 160 | WRITE | 50 | 2127 | 0 | 308 | 72.0% |

■ `printa(ioStat, "%d KB | %8s | %6x %-6d %06d %6d")`

| intKey0 | KB | | strKey0 | | avg | count | min | max | pct% |
|---------|----|--|---------|--|------|-------|--------|-----|-------|
| 96 | KB | | WRITE | | 0x1e | 46 | 000000 | 210 | 0.9% |
| 8 | KB | | READ | | 0x17 | 69 | 000000 | 175 | 1.1% |
| 12 | KB | | READ | | 0x19 | 65 | 000000 | 218 | 1.1% |
| 128 | KB | | READ | | 0x5 | 2482 | 000000 | 727 | 8.8% |
| 160 | KB | | WRITE | | 0x32 | 2127 | 000000 | 308 | 72.0% |

Formatting Aggregate Output – Examples Contd.

- `printa(ioStat, "%d KB | %8s | %avg$d | %count$d")`

| intKey0 KB | strKey0 | avg | count |
|------------|---------|-----|-------|
| 96 KB | WRITE | 30 | 46 |
| 8 KB | READ | 23 | 69 |
| 12 KB | READ | 25 | 65 |
| 128 KB | READ | 5 | 2482 |
| 160 KB | WRITE | 50 | 2127 |

- `printa(ioStats, "{ len: %=d, type: \"%=s\", avg: %avg$d, count: %count$d },")`

```
{ len: 96, type: "WRITE", avg: 30, count: 46 },  
{ len: 8, type: "READ", avg: 23, count: 69 },  
{ len: 12, type: "READ", avg: 25, count: 65 },  
{ len: 128, type: "READ", avg: 5, count: 2482 },  
{ len: 160, type: "WRITE", avg: 50, count: 2127 },
```

Shared Integers

- New storage classes for integers:
 - `pervm` – shared in a VM, across all VCPU and VMX threads
 - `pervmk` – shared in the VMkernel, across all PCPUs
 - `perhost` – shared in the ESX host, across all VMs and the Vmkernel
- New built-ins `fetchadd` and `cmpxchg`
- Reads and writes are fenced for sequentially consistency

```
pervm int hvExitCtr;

HVExit {
    fetchadd(hvExitCtr, 1);
}

VMXUnload {
    printf("%u\n", hvExitCtr);
}
```

Shared Integers – Example

```
// excerpt from --  
// bora/devkits/tools/ddv/scripts/ddv_common.emt  
pervmk int bigDDVLock;  
  
int ddvSpinLockAcquire()  
{  
    int lockAcquired;  
    if (!cmpxchg(bigDDVLock, 0, 1)) {  
        lockAcquired = ddvSpinLockRetry(1);  
    } else {  
        lockAcquired = SPLOCK_SUCCESS;  
    }  
    return lockAcquired;  
}
```


Emmett Pre-Processor (EPP)

- Much like the C pre-processor
- **Macros:** `#define`, `#undef`
 - `-D IDENT [=VAL]` vprobe app switch
 - No support for macro arguments
- **File inclusion:** `#include`
 - `-I PATH` vprobe app switch
- **Conditional inclusion:** `#if`, `#ifdef`, `#elif`, `#else`
- **Misc:** `#error`, `#warning`, `#line`
- **Pre-defined macros:** `__VERSION__`, `__ESX__`, `__DESKTOP__`
- Part of Emmett compiler, written in OCaml.
 - CPP doesn't work due to distribution issues (Windows, OSX, Linux, ESX)

EPP – Example

```
// excerpt from
// bora/devkits/tools/ddv/scripts/ddv_common.emt

#if (__VERSION__ < 2013)
int ddvCurModID()
{
    ModInfoStack *modInfoStack;
    modInfoStack = curworld()->modStack;
    return (modInfoStack ? modInfoStack->modID : DUT_MOD_ID_INVALID);
}

#define CUR_MOD_ID ddvCurModID()
#else
#define CUR_MOD_ID LASTMODID
#endif /* (__VERSION__ < 2013) */
```

Probe Globbing

- Define probe points using regex patterns
- Available in the VMK domain

- Example:

```
aggr histo[0][1];  
VMK:ENTER:Timer_* { histo[PROBENAME]++; }  
VMK:VMKUload      { printa(histo); }
```

- Output:

| strKey0 | count | pct% |
|-------------------------------------|--------|-------|
| VMK:ENTER:Timer_WorldPreCleanup | 1 | 0.0% |
| VMK:ENTER:Timer_RemoveSync | 33 | 0.0% |
| VMK:ENTER:Timer_ModifyOrAddTC | 60 | 0.0% |
| VMK:ENTER:Timer_SysUptimeUS | 63 | 0.0% |
| VMK:ENTER:Timer_SysUptime | 121 | 0.0% |
| VMK:ENTER:Timer_UpdateOneShot | 312 | 0.1% |
| VMK:ENTER:Timer_Pending | 362 | 0.1% |
| VMK:ENTER:Timer_Stats | 705 | 0.3% |
| VMK:ENTER:Timer_Remove | 813 | 0.4% |
| VMK:ENTER:Timer_AddTC | 917 | 0.4% |
| VMK:ENTER:Timer_SignedTCToMS | 954 | 0.5% |
| VMK:ENTER:Timer_AddTCWithLockDomain | 1279 | 0.6% |
| VMK:ENTER:Timer_EnableStatsTimer | 3560 | 1.8% |
| VMK:ENTER:Timer_Interrupt | 5884 | 3.1% |
| VMK:ENTER:Timer_BHHandler | 26241 | 13.9% |
| VMK:ENTER:Timer_GetCycles | 146500 | 78.0% |

Fault Injection

Fault Injection Framework

- Destructive builtins that modify system state
- `setphysgpr()`
 - Modify physical register state
 - Available in the VMM and VMK domains
- `setgpr()`
 - Modify guest register state (i.e., virtual registers)
 - Available in the VMM and GUEST domains
- `setvmw()`
 - Modifies memory contents
 - Available in the VMK domain
- `genmce()`
 - Inject MCE's into the guest

Fault Injection Framework

- `skipfunction()`
 - Modifies control flow
 - Skip execution of the current function
 - Available in VMK, at function entry points
- Example

```
VMK:ENTER:vmklinux_9.vmklnx_kmalloc {  
    if (injectFault) {  
        /* set return value as NULL */  
        setphysgpr(REG_RAX, 0);  
        /* skip execution of vmklnx_kmalloc*/  
        skipfunction();  
    }  
}
```

VMK Watchpoints

- Program physical DR's for read/write watchpoints
- Builtins: `setwatchpoint()`, `removewatchpoint()`
- Static probe: WatchpointHandler
 - Fire probe when the watchpoint triggers
- Limitations
 - Use NMIs to propagate DR updates to other PCPUs
 - VMM is not NMI hardened
 - Hence, this is an unsupported feature

VMK Watchpoints Example

```
pervmk int rsp, wpid;
ENTER:Timer_Interrupt {
    if (PCPU == 0) {
        rsp = getphysgpr(REG_RSP);
        wpid = setwatchpoint(rsp, 0x2/*8 bytes*/, 0x1/*WRITE*/);
    }
}

WatchpointHandler(int addr) {
    string bt;
    vmwstack(bt);
    printf("Write watchpoint at %p, backtrace %s\n", addr, bt);
}

EXIT:Timer_Interrupt {
    if (PCPU == 0) {
        removewatchpoint(rsp);
    }
}
```


Usability Enhancements

External Types

- Let's write a script that prints `PRDA->lastIntIdx` every second, for every PCPU
 - You must manually determine `lastIntIdx`'s offset and declare a sparse type

```
typedef struct PRDA {  
    @0x54 lastIntIdx;  
} PRDA;  
  
VMK1Hz {  
    PRDA *p = PRDA_GET_ADDR;  
    printf("PCPU %d: %d\n", PCPU, p->lastIntIdx);  
}
```

External Types

- New external type support allows referencing VMK types without the need for declaration:

```
VMK1Hz {  
    $vmkernel.PRDA *p = PRDA_GET_ADDR;  
    printf("%d\n", p->lastIntIdx);  
}
```

- Notation: *\$MODULE.TYPE*

- \$vmkernel.World_Handle
- struct \$vsan.BatchResult
- union \$vmkernel.PCIENCapReg
- enum \$tcpip4.VmkNetDomain = \$tcpip4.VMK_INET_DOMAIN

Line Numbers

- Let's write a script that probes a line of function World_New() and prints the PCPU#

```
/* bora/vmkernel/main/world.c */  
3012: VMK_ReturnStatus  
3013: World_New(World_InitArgs *args,  
3014:           World_ID      *worldID)  
    ...  
3035: if (args->typeFlags & WORLD_USER) {  
3036:     initTable = userTableInit;  
    ...
```

- You must manually map world.c:3036 to the corresponding function-relative byte offset

```
VMK:OFFSET:World_New:0xA8 {  
    printf("PCPU %d\n", PCPU);  
}
```

Line Numbers

- New line number probing support eliminates need for manual line number to byte offset translation
- File-relative line numbers:

```
VMK:OFFSET:World_New:L3036 {  
    printf("PCPU: %d\n", PCPU);  
}
```

- Function-relative line numbers:

```
VMK:OFFSET:World_New:F23 {  
    printf("PCPU: %d\n", PCPU);  
}
```

Local Variables and Arguments

- Now we want to inspect `args->p2mCacheSize` at the same probe point (`world.c:3036`)

```
/* bora/vmkernel/main/world.c */
3012: VMK_ReturnStatus
3013: World_New(World_InitArgs *args,
3014:           World_ID      *worldID)
    ...
3035: if (args->typeFlags & WORLD_USER) {
3036:     initTable = userTableInit;
    ...
```

Local Variables and Arguments

- Must manually determine:
 - the offset of `p2mCacheSize` in `World_InitArgs`
 - where `args` is stored at the probe point (stack? register?)

```
typedef struct World_InitArgs {  
    @0x80 uint32 p2mCacheSize;  
} World_InitArgs;  
  
VMK:OFFSET:World_New:0xA8 {  
    World_InitArgs *args = getphysgpr( REG_RBX );  
    printf("%d\n", args->p2mCacheSize);  
}
```

Local Variables and Arguments


- New local variable inspection support allows target local variables to be directly referenced by name:

```
VMK:OFFSET:World_New:L3036 {  
    printf("%d\n", $args->p2mCacheSize);  
}
```


- emmett translates the name reference to the correct register/memory lookup under-the-hood
- Arguments and local variables are treated the same way
- No support yet for static locals

Symbol Databases

- Pre-requisite to external types, line numbers and local variable support
- Symbol database creation (once per build):
 - `bora/vmcore/support/vprobes/util/createEmtSymbolDb.py`



- Symbol database usage:



Useful Links

- **User Guide:**

<http://engweb.eng.vmware.com/monitor/vprobes/doc/html/userGuide.html>

- **Mailing list**

vprobes@vmware.com

- **Web site**

<http://vprobes.eng.vmware.com>