

How to identify linux kernel issue?

Kernel Panic

Kernel Hang

Performance issue

Kernel Panic

what's kernel panic

```
/**
 * panic - halt the system
 * @fmt: The text string to print
 *
 * Display a message, then perform cleanups.
 *
 * This function never returns.
 */
void panic(const char *fmt, ...)
{
```

Linux Guest OS problems from customer

1. Guest OS rebooting unexpectedly

Possible reasons:

1. Rebooted manually. <---/var/log/message
2. Guest OS panic and kdump enabled. <---/var/crash/vmcore

3. Guest OS hung

Possible reasons:

1. Guest OS kernel panics and stops there. <-----Do not reboot Guest OS, collect vmss
2. Guest OS hung. <-----Do not reboot Guest OS, collect vmss

3. Low performance, debug inside guest OS

What is a vmcore?

A vmcore (Virtual Memory Core) is a dump of all memory contents on a system at a point in time.
Can be debugged by crash tool

Why might I want to capture a vmcore?

Vmcores can be useful in many situations:

- They capture system state at a particular point in time. This is often useful at the time of a crash, or at a user-defined time when crashed manually by the customer.
- Can be used to see why a system might be slow. Can analyse what the system is doing at the time.
- <there are many more reasons>

How do I capture a vmcore?

Config kdump and capture vmcore.

Suspend or take a snapshot on VM, collect vmss/vmsn

1. Kdump

- 1.reserve physical memory for second kernel
2. run kexec system call to start second kernel
- 3.collect vmcore by second kernel.

2. Suspend or take a snapshot on VM, collect vmss/vmsn

#vmss2core -N6 xxx.vmss

If RHEL6 kernel, to reduce the vmss.core size.

```
#makedumpfile -f -d 31 -x vmlinux vmss.core vmcore
```

How to open a vmcore?

1. find kernel version:

```
#strings vmcore|grep "2.6."
```

2. extrace debuginfo rpm

```
#rpm2cpio kernel-debuginfo-2.6.18-128.el5.x86_64.rpm | cpio -idv
```

3. crash vmlinux vmcore

Environment

host: cybertron.eng.vmware.com

user/passwd: gss/vmware

1. Kernel Panic

kernel Panic analysis

1. Find kernel oops in the log

```
crash> log
```

2. Print backtrace of panic thread.

```
crash> bt
```

Other commands:

ps - list processes

set <pid> switch process

set -p switch back to panic process

sym - source code file and line number

dis - disassemble the address

rd - read memory

2. Kernel Hang

Hung Caused by D state process

1. First find UNINTERRUPTABLE process

```
ps -ef|grep UN
```

2. switch to UN process and print call stack

```
set <pid>
```

```
bt or foreach UN bt
```

3. Find the lock owner by reading stack data.

```
rd xxxxx -e xxxxx -S
```

file dentry...

Hung caused by CPU softlock up

```
#sysctl -w kernel.softlockup_panic=1
```

1. find soft lockup process in dmesg

```
#log
```

2. switch to softlock up process

```
#set <pid>
```

3. print stack

```
#bt
```

Others

Out of Memory

```
#kmem -i
```

reached max pid

```
#ps -ef|wc -l
```

3.Low performance, debug inside guest OS

3.1 sysrq.

enable sysrq

```
sysctl -w kernel.sysrq=1
```

1.to panic guest OS

```
echo c >/proc/sysrq-trigger
```

2.dump stacks

```
echo t >/proc/sysrq-trigger
```

3.dump memory

```
echo m >/proc/sysrq-trigger
```

3.2. perf

```
#perf top
```

```
#perf top -p <pid>
```

3.3 vmcore

```
#runq
```

```
#foreach RU bt
```

```
#ps|grep RU
```

Useful links

Linux coredump analyzer:

<http://cybertron.eng.vmware.com>

LCSA:

<http://lcsa.eng.vmware.com:5001/>

kernel oops analyzer:

<https://access.redhat.com/labs/kerneloopsanalyzer>

rhel kernel source code:

<https://access.redhat.com/lab/psb>