


Memory Management

1. Numa

what's a numa?

Non-Uniform Memory Access (NUMA) refers to multiprocessor systems whose memory is divided into multiple memory nodes. The access time of a memory node depends on the relative locations of the accessing CPU and the accessed node.



How to check numa status?

```
[root@pek2-office-18th-10-117-174-219 ~]# numactl --hardware
```

available: 1 nodes (0)

node 0 cpus: 0 1 2 3 4 5 6 7

node 0 size: 15994 MB

node 0 free: 226 MB

node distances:

node 0

0: 10

Numa migration

Move pages from one node to another. system call is provided by kernel.

NAME

migrate_pages - move all pages in a process to another set of nodes

SYNOPSIS

```
#include <numaif.h>
```

```
long migrate_pages(int pid, unsigned long maxnode,
                  const unsigned long *old_nodes,
                  const unsigned long *new_nodes);
```

https://www.kernel.org/doc/Documentation/vm/page_migration

2. Zone

What's a zone?

Kernel divided physical memory with in a numa node into zones based on various restrictions on how it can be used.

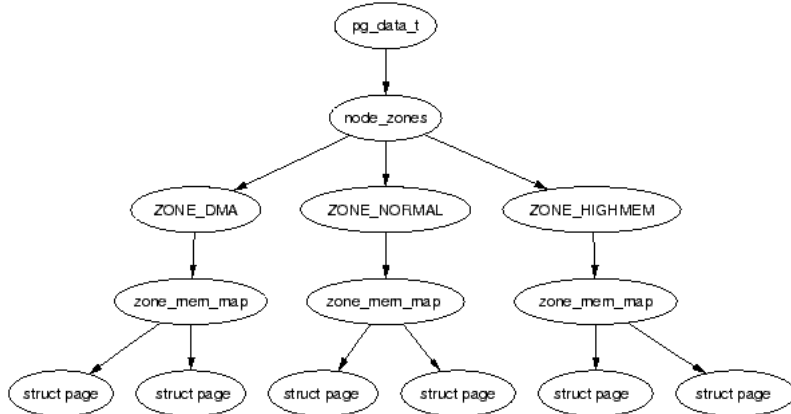


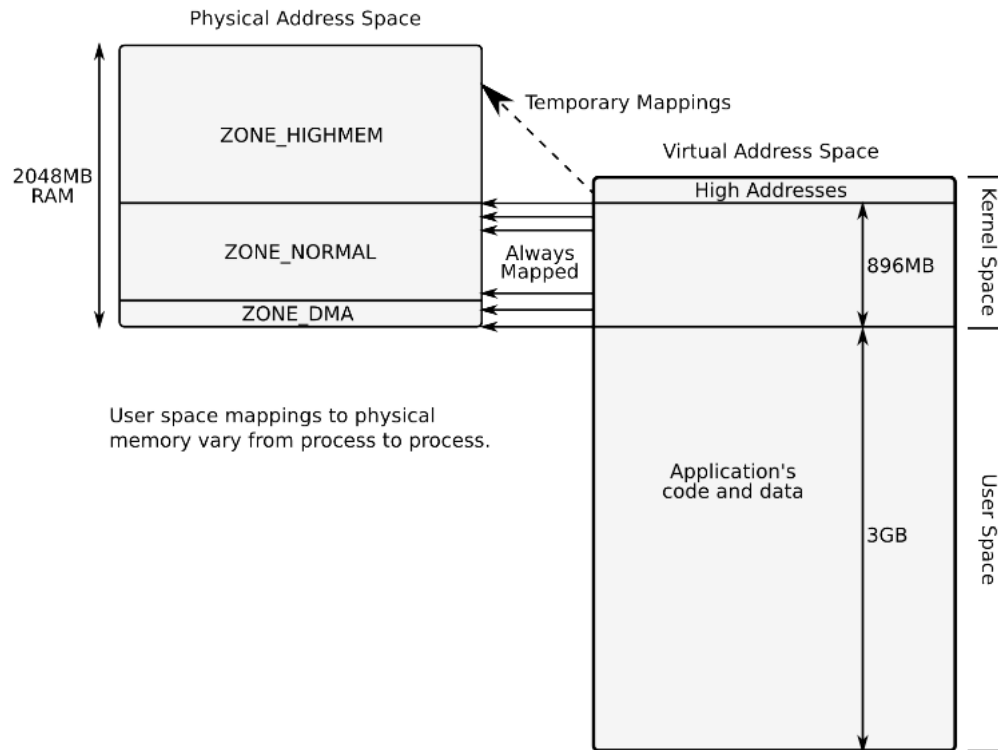
Figure 2.1: Relationship Between Nodes, Zones and Pages

How about Android kernel?

Yes, it has only normal and high zone!!!

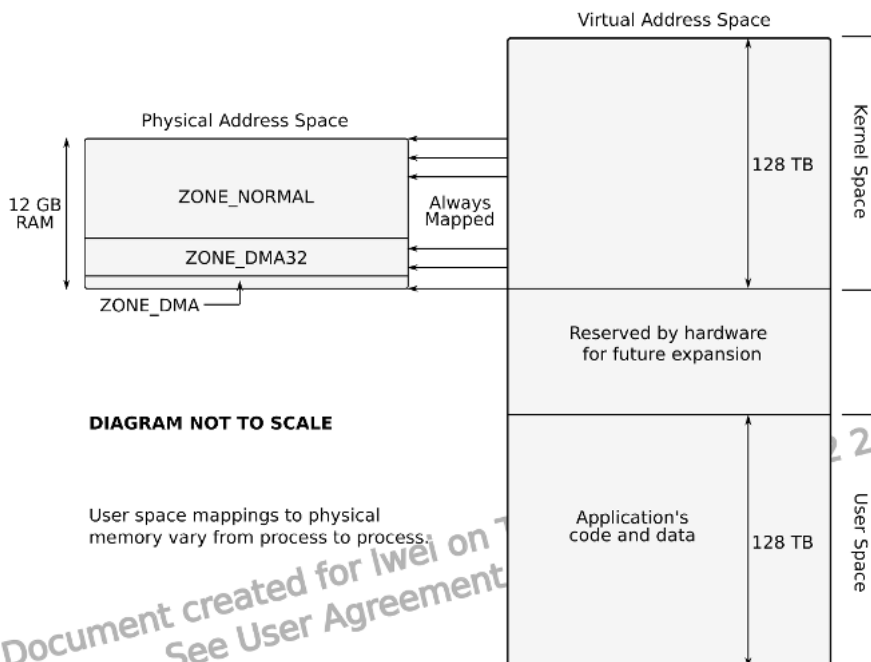
```
127|u0_a150@android:/ $ cat /proc/buddyinfo
Node 0, zone Normal 7506 1472 118 16 12
Node 0, zone HighMem 338 41 1 0 3
u0_a150@android:/ $
```

X86 32 address space mapping




X86-64 address space mapping

Figure 10.1. x86-64 virtual address space



<http://www.ibm.com/developerworks/cn/linux/l-numa/>

watermark of a zone**Emergent Pool**

```
#define ALLOC_HIGH 0x20 /* __GFP_HIGH set */
if (alloc_flags & ALLOC_HIGH)
min -= min / 2;
```

low mem reserve

low mem reserve is used protect the low memory zone by raising watermark up when there's failure in up level zone.

How to check check the workmark of zone?

```
[gliu@pek2-office-18th-10-117-174-219 ~]$ cat /proc/zoneinfo
```

```
Node 0, zone  DMA
pages free  3871
  min      16
  low      20
  high     24
 scanned   0
 spanned   4095
 present   3998
 managed   3977
```

or

```
echo m >/proc/sysrq-trigger ; tail /var/log/message -n 100
```

```
crash> kmem -n
```

```
NODE SIZE  PGLIST_DATA  BOOTMEM_DATA  NODE_ZONES
 0 524285 ffff88000000a000 ffffffff81cae100 ffff88000000a000
                                ffff88000000126c0
                                ffff880000001ad80
                                ffff8800000023440
MEM_MAP  START_PADDR  START_MAPNR
ffffea0000000038 1000 1
```

```
ZONE NAME  SIZE  MEM_MAP  START_PADDR  START_MAPNR
```

0	DMA	4095	ffffea0000000038	1000	1
1	DMA32	520190	ffffea0000038000	1000000	4096
2	Normal	0	0	0	0
3	Movable	0	0	0	0

```
crash> kmem -z
```

```
NODE: 0 ZONE: 0 ADDR: ffff8800000a000 NAME: "DMA"
SIZE: 4095 PRESENT: 3841 MIN/LOW/HIGH: 83/103/124
VM_STAT:
NR_FREE_PAGES: 3938
```

```
NODE: 0 ZONE: 1 ADDR: ffff8800000126c0 NAME: "DMA32"
SIZE: 520190 PRESENT: 513078 MIN/LOW/HIGH: 11180/13975/16770
VM_STAT:
NR_FREE_PAGES: 270234
```

Tuning of min watermark and low memory reserve?

```
vm.min_free_kbytes
```

```
vm.lowmem_reserve_ratio
```

3.page

data struct of a page

```
struct page {
    long unsigned int flags; <---page flags, like lru/dirty/writeback/active/inactive/slab/highmem...
    atomic_t _count; <----how many processes uses this page?
                                checked this counter when dropping a cache.

    union {
        atomic_t _mapcount; <---how many page table(pte) points to this page
        struct {
            u16 inuse;
            u16 objects;
        };
    };
    union {
        struct {
            long unsigned int private;
            struct address_space *mapping; <-----if this page is file backed, it points to the address_space,
                                            if it's anon page, points anon_struct
        };
        spinlock_t ptl;
        struct kmem_cache *slab;
        struct page *first_page;
    };
    union {
        long unsigned int index; <-----if this page if file backed, index is used to store the offset in this file,
                                mean the position of this page within a file.

        void *freelist;
    };
    struct list_head lru; <----used for PFRA
}
```

mem_map

This is a array to store all the physical page infomation.

zone has a index to store where this zone starts at mem_map

```
crash> whatis zone.zone_start_pfn
```

```
struct zone {
    [34336] long unsigned int zone_start_pfn;
}
```

mem_map array take about 1.36% of total memory
mem_map ratio = sizeof(page)/4096

4. buddy allocator

what's buddy allocator?




Figure 6.1: Free page block management

buddy allocator interface

struct page * alloc_page(unsigned int gfp_mask)
◆Allocate a single page and return a struct address
struct page * alloc_pages(unsigned int gfp_mask, unsigned int order)
◆Allocate 2 ^{order} number of pages and returns a struct page
unsigned long get_free_page(unsigned int gfp_mask)
◆Allocate a single page, zero it and return a virtual address
unsigned long __get_free_page(unsigned int gfp_mask)
◆Allocate a single page and return a virtual address
unsigned long __get_free_pages(unsigned int gfp_mask, unsigned int order)
◆Allocate 2 ^{order} number of pages and return a virtual address
struct page * __get_dma_pages(unsigned int gfp_mask, unsigned int order)
◆Allocate 2 ^{order} number of pages from the DMA zone and return a struct page

GFP flag

Flag	Description
__GFP_WAIT	Indicates that the caller is not high priority and can sleep or reschedule
__GFP_HIGH	Used by a high priority or kernel process. Kernel 2.2.x used it to determine if a process could access emergency pools of memory. In 2.4.x kernels, it does not appear to be used
__GFP_IO	Indicates that the caller can perform low level IO. In 2.4.x, the main affect this has is determining if try_to_free_buffers() can flush buffers or not. It is used by at least one journaled filesystem
__GFP_HIGHIO	Determines that IO can be performed on pages mapped in high memory. Only used in try_to_free_buffers()
__GFP_FS	Indicates if the caller can make calls to the filesystem layer. This is used when the caller is filesystem related, the buffer cache for instance, and wants to avoid recursively calling itself

Flag	Low Level Flag Combination
GFP_ATOMIC	HIGH
GFP_NOIO	HIGH WAIT
GFP_NOHIGHIO	HIGH WAIT IO
GFP_NOFS	HIGH WAIT IO HIGHIO

GFP_KERNEL	HIGH WAIT IO HIGHIO FS
GFP_NFS	HIGH WAIT IO HIGHIO FS
GFP_USER	WAIT IO HIGHIO FS
GFP_HIGHUSER	WAIT IO HIGHIO FS HIGHMEM
GFP_KSWAPD	WAIT IO HIGHIO FS

Avoid memory fragmentation???

1. memory compaction

<http://lwn.net/Articles/368869/>

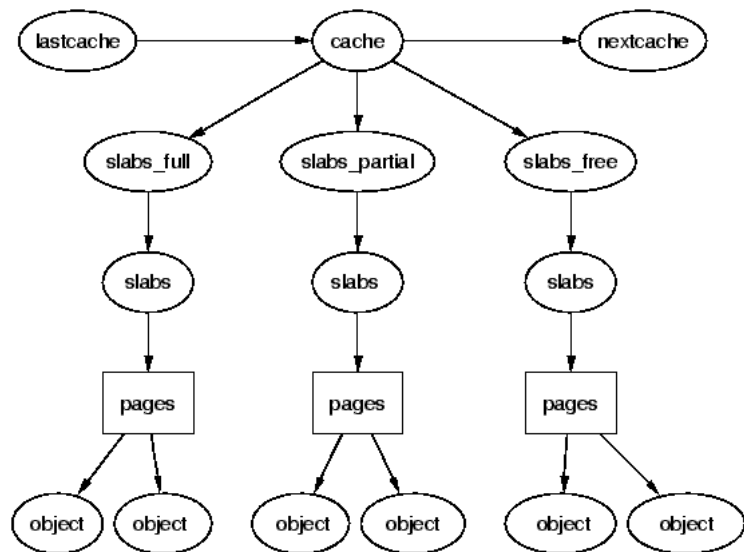
2. zone_movable

<http://lwn.net/Articles/224255/>

3. lumpy reclaim

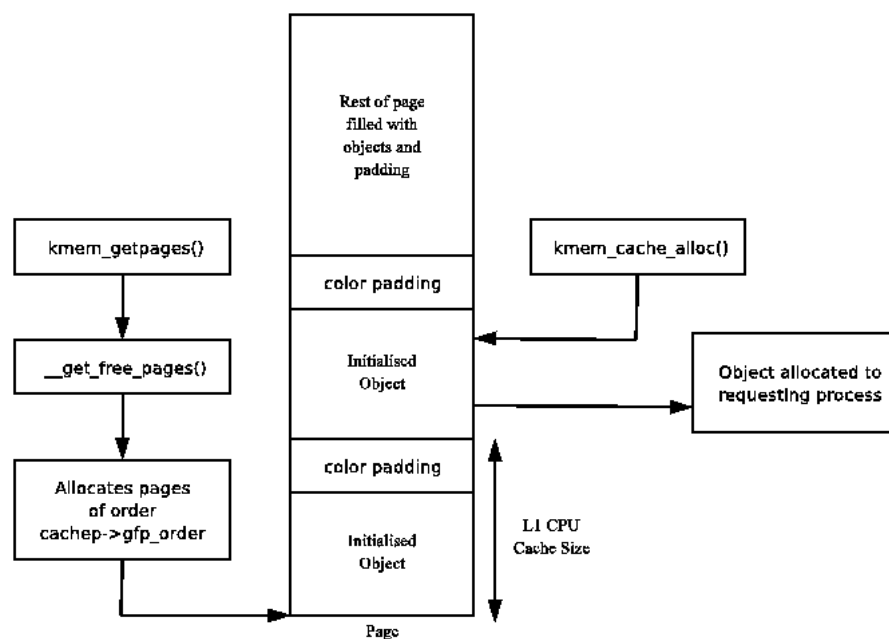
<http://lwn.net/Articles/211199/>

5. slab allocator

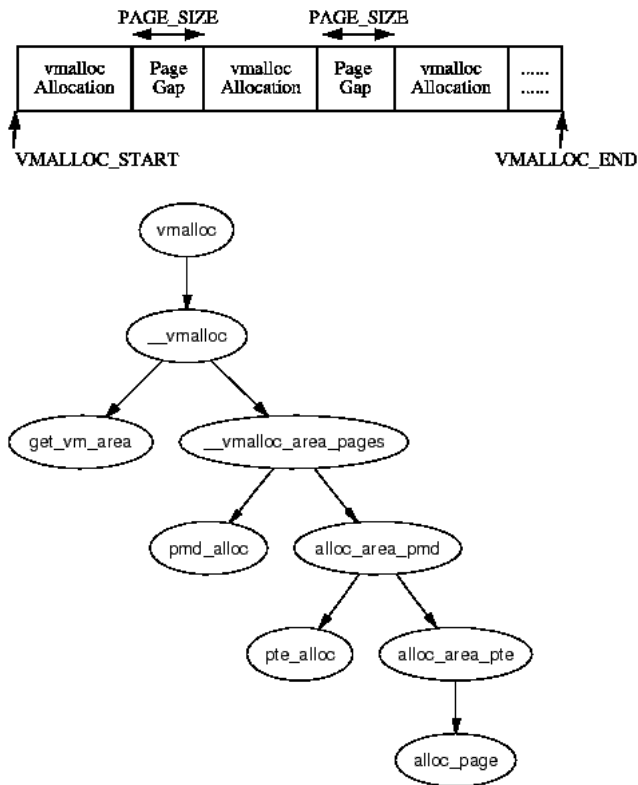


Color of slab?

In this scenario, the first slab created will have its objects start at 0. The second will start at 32, the third at 64, the fourth at 96 and the fifth will start back at 0. With this, objects from each of the slabs will not hit the same hardware cache line on the CPU. The value of `colour` is 3 and `colour_off` is 32.



6. vmalloc



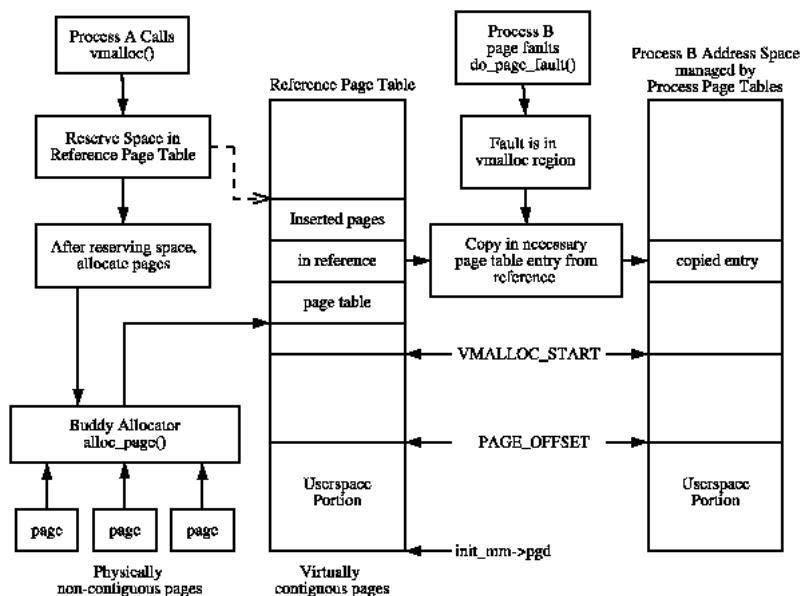
Why vmalloc can not be used in interrupt routine?

1. vmalloc will call kcalloc(GFP_KERNEL) to allocate page table, which will cause process to sleep.

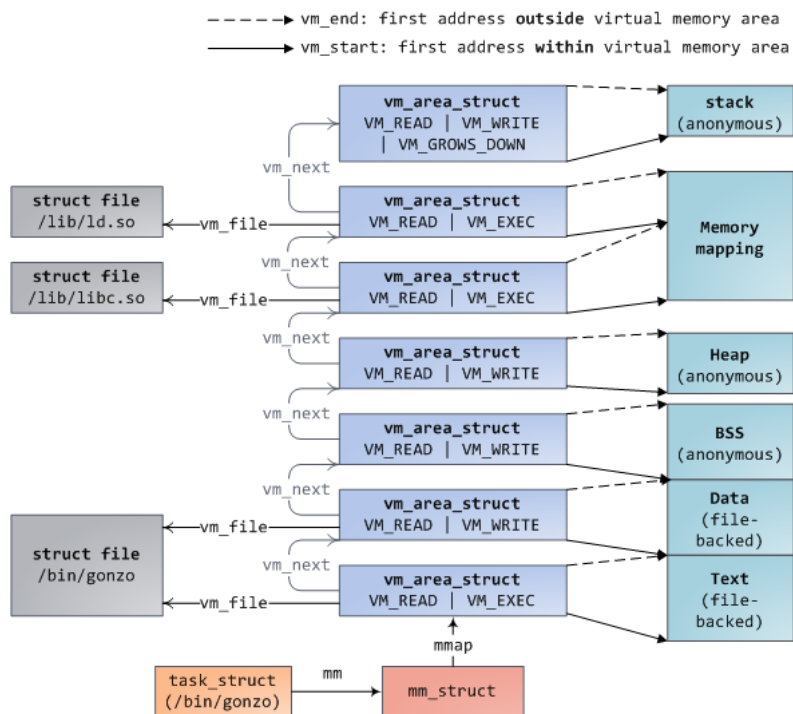
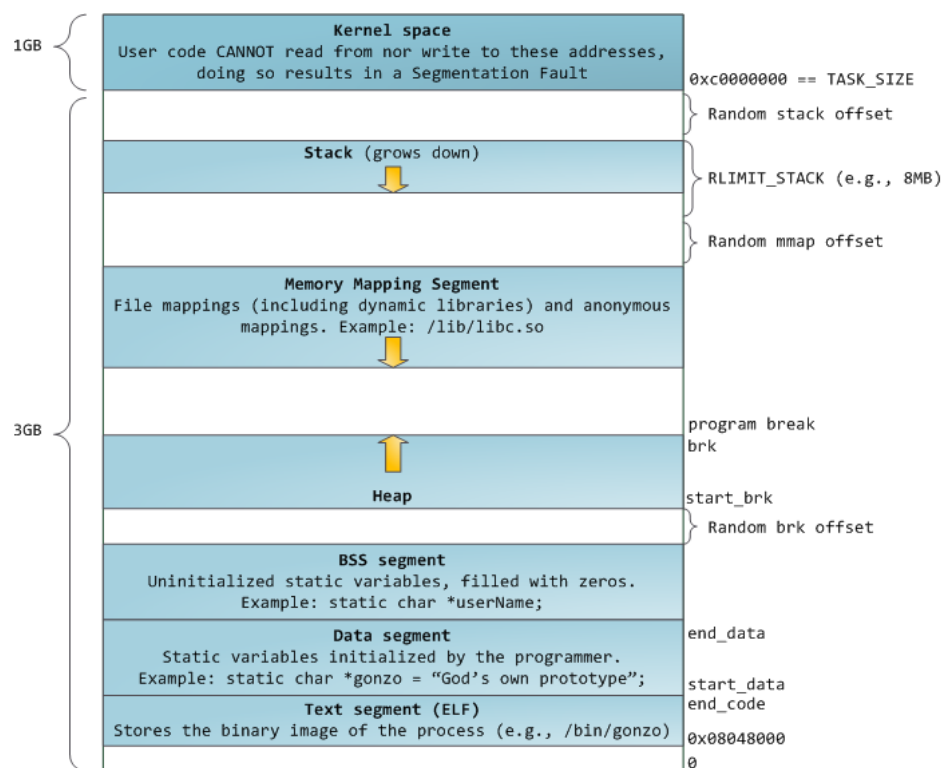
2. vmalloc will cause page fault

The page table updated by vmalloc() is not the current process but the reference page table stored at init_mm->pgd.

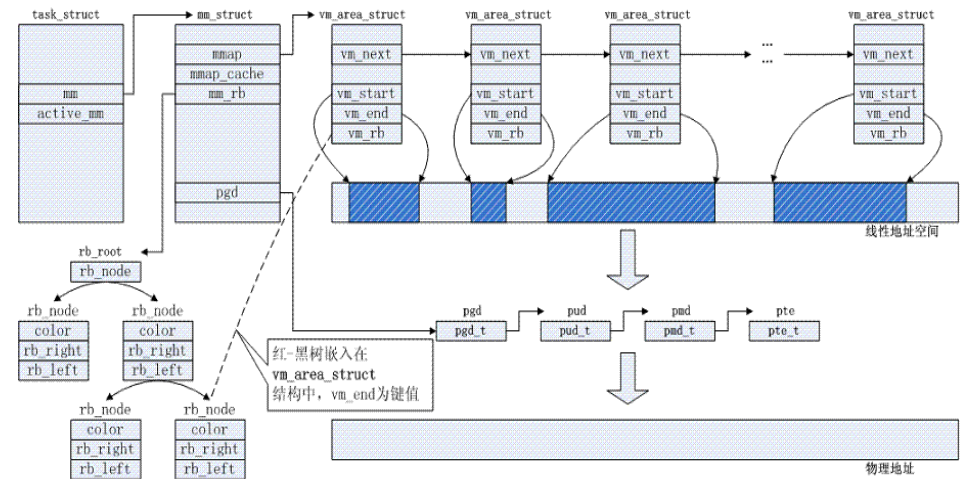
This means that a process accessing the vmalloc area will cause a page fault exception as its page tables are not pointing to the correct area. There is a special case in the page fault handling code which knows that the fault occurred in the vmalloc area and updates the current process page tables using information from the master page table.



7. Process Address space



vm_area_struct



<http://duartes.org/gustavo/blog/post/how-the-kernel-manages-your-memory/>

How to check the vm_area_struct

```
[root@pek2-office-18th-10-117-174-219 work]# pmap -x 11749
11749: ./loop
```

Address	Kbytes	RSS	Dirty	Mode	Mapping
0000000000400000	4	4	0	r-x--	loop
0000000000600000	4	4	4	r----	loop
0000000000601000	4	4	4	rw---	loop
0000003d91200000	132	128	0	r-x--	ld-2.17.so
0000003d91420000	4	4	4	r----	ld-2.17.so
0000003d91421000	4	4	4	rw---	ld-2.17.so
0000003d91422000	4	4	4	rw---	[anon]
0000003d91600000	1752	640	0	r-x--	libc-2.17.so
0000003d917b6000	2048	0	0	-----	libc-2.17.so
0000003d919b6000	16	16	8	r----	libc-2.17.so
0000003d919ba000	8	8	8	rw---	libc-2.17.so
0000003d919bc000	20	12	12	rw---	[anon]
00007f19d5d6d000	12	12	12	rw---	[anon]
00007f19d5d90000	4	4	4	rw---	[anon]
00007fff41214000	132	12	12	rw---	[stack]
00007fff4124f000	8	0	0	r----	[anon]
00007fff41251000	8	4	0	r-x--	[anon]
ffffffff600000	4	0	0	r-x--	[anon]

```
crash> vm
PID: 2596 TASK: ffff88007b0e1540 CPU: 0 COMMAND: "crash"
MM PGD RSS TOTAL_VM
ffff88007c4f2d00 ffff88007d4e5000 828332k 972200k
VMA START END FLAGS FILE
ffff88007d4e3530 400000 a1e000 8001875 /usr/bin/crash
ffff88007d4e35f8 c1d000 c3f000 8101873 /usr/bin/crash
ffff88007d4e36c0 c3f000 dd5000 100073
ffff88007d4e3788 e3e000 e5e000 8101873 /usr/bin/crash
ffff880078935ae8 17f3000 298c8000 100073
ffff88007c36d3e0 302ae00000 302ae02000 8000075 /lib64/libdl-2.12.so
ffff88007c36d318 302ae02000 302b002000 8000070 /lib64/libdl-2.12.so
ffff88007c36d4a8 302b002000 302b003000 8100071 /lib64/libdl-2.12.so
```

mprotect

NAME

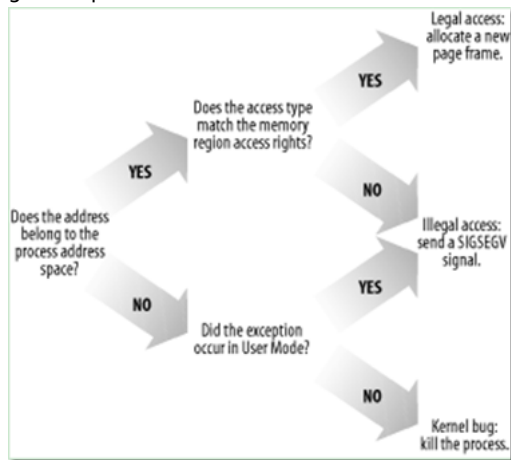
mprotect - set protection on a region of memory

SYNOPSIS

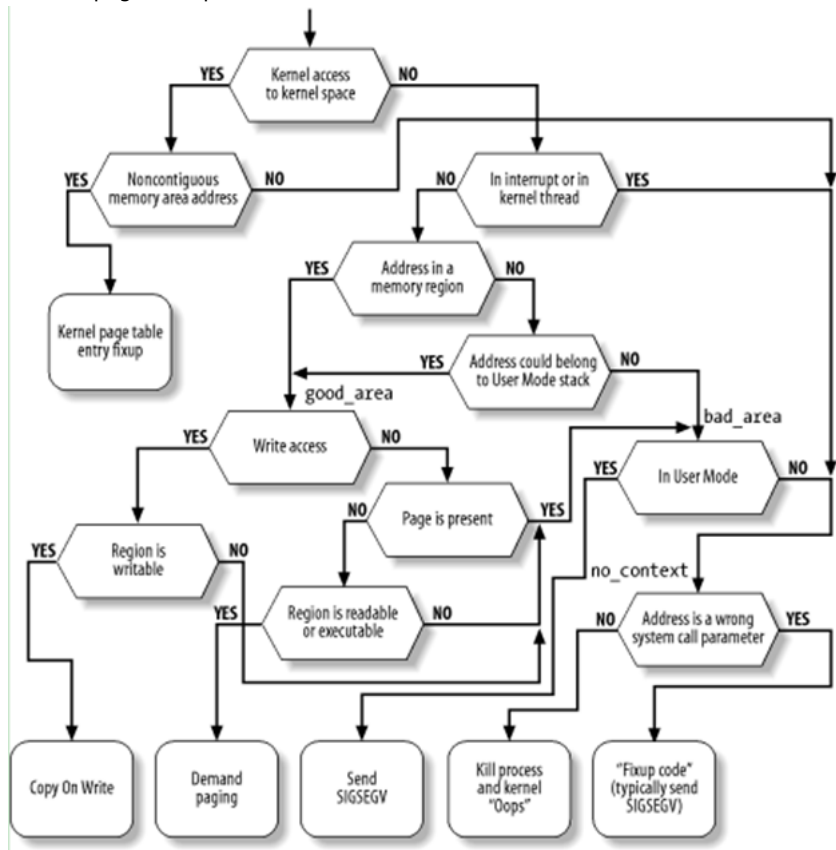
```
#include <sys/mman.h>
int mprotect(void *addr, size_t len, int prot);
```

Page Fault


general procedure



detailed page fault procedure



Page on demand



new and malloc()


how about this code?

```
int *p = malloc(1024*1024*sizeof(int));
for(i = 0; i < 4096*1024;i++){
```

```
printf("%d", p[i]);
}
```

In this case, the system-wide `empty_zero_page`, which is just a page of zeros, is mapped for the PTE and the PTE is write protected.

COW




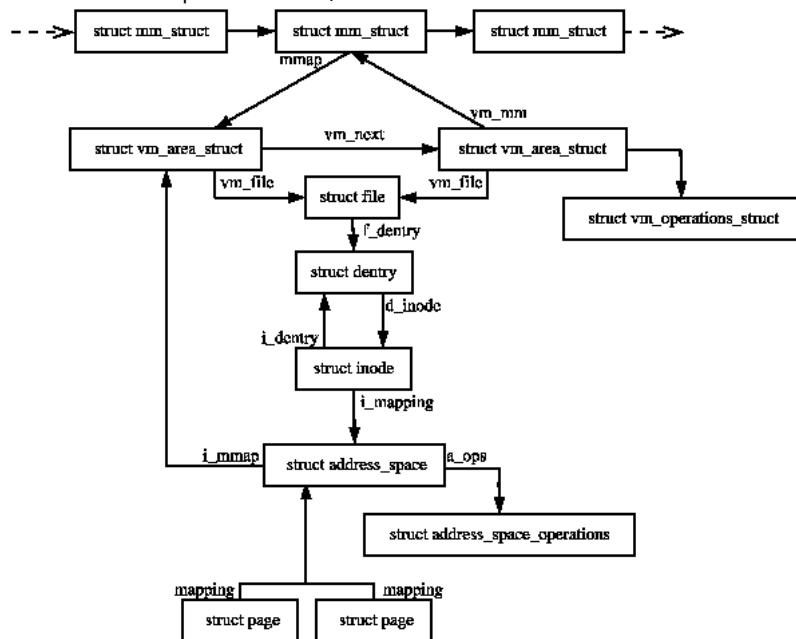
Linux recognises a COW page because even though the PTE is write protected, the controlling VMA shows the region is writable.

8. Memory Cache/Buffer


address_space

what's a address_space?

a structure to represent cache/buffer



radix tree



radix tree tagging

mark a rb-node and all its parent recursively until root if the one of its page if dirty or writeback

share memory

tmpfs based
mapped the tmpfs file to process's address space

<https://www.kernel.org/doc/gorman/html/understand/understand015.html>

9. PFRA

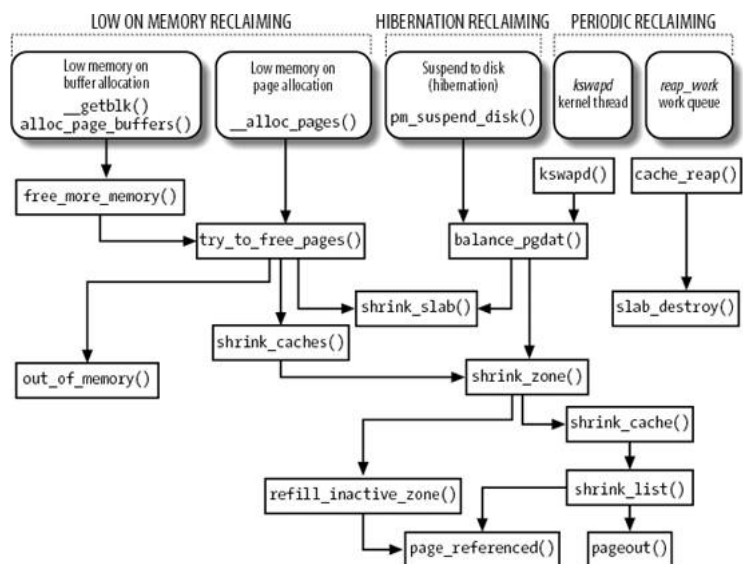
which page to be reclaimed?

Type of pages	Description	Reclaim action
Unreclaimable	Free pages (included in buddy system lists)	(No reclaiming allowed or needed)
	Reserved pages (with <code>PG_reserved</code> flag set)	
	Pages dynamically allocated by the kernel	
	Pages in the Kernel Mode stacks of the processes	
	Temporarily locked pages (with <code>PG_locked</code> flag set)	
Swappable	Memory locked pages (in memory regions with <code>VM_LOCKED</code> flag set)	Save the page contents in a swap area
	Anonymous pages in User Mode address spaces	
Syncable	Mapped pages of <code>tmpfs</code> filesystem (e.g., pages of IPC shared memory)	Synchronize the page with its image on disk, if necessary
	Mapped pages in User Mode address spaces	
	Pages included in the page cache and containing data of disk files	
	Block device buffer pages	
Discardable	Pages of some disk caches (e.g., the inode cache)	Nothing to be done
	Unused pages included in memory caches (e.g., slab allocator caches)	
	Unused pages of the dentry cache	

Algorithm??

second chance and lru combined
active referenced

How to reclaim?



10. SWAP

to be continued...