

XvMotion Internals

Arun

March 11, 2013

Confidential

vmware®

Agenda

XvMotion buffer management

- Migrate module heap overview
- Per Migration heap
- XvMotion buffer - Bitmaps, bloomFilter, disk queues and block buffers

XvMotion storage stream

- Storage stream generator
- Storage stream handler

XvMotion IO

- Source - clone and mirror block allocation and read
- Destination - handle conflicting IO/Async IO write

XvMotion Buffer Management

Overview migrate module heaps

Migrate Module heap

- Dynamic heap
- Create at the module load time
- Small
- Mainly for meta data allocation
 - Per migration slice and XvMotion slice meta data
- Different from per migration default heap!

init_module

Migrate_Init

MigrateModuleHeapInit

Heap_DetermineMaxSize (Migration Slices and XvMotion Slices)

Heap_Create

Per Migration Default Heap

- Static Heap
- Default heap is create at module load time
- If default heap is not available then heap is create during InitMigration time

init_module

 Migrate_Init

 PagePool_Create

 MigrateVASpaceInit

 VASpace_Reserve (12 * 768 MB)

 MigrateHeapCreateDefault

 MigrateHeapDetermineDefaultSize

 Migrate_HeapCreate

Per Migration Heap Creation

- If default heap is not available then create a heap

Migrate_InitMigration

MigrateInfo_Alloc

Migrate_HeapCreate

MigrateVASpaceCreateHeap

MigrateVASpaceReserveSlice

For numpages

Migrate_NiceAlloc

PagePool_Alloc

VASpace_Map

Heap_CreateStatic

XvMotion Buffer Layout

XvMotion buffer layout

```
/*
 * Each VASpace slice is laid out in the following structure:
 *
 * +-----+ 0 + vpnoffset
 * |   Page Bitmap   | 1 to 2 Pages
 * +-----+
 * |   Block Bitmap  | 4 to 16 Pages
 * +-----+
 * | Per-device Bloom Filter | 2 pages per device, 276 pages
 * +-----+
 * |   Queue          | 128 to 512 Pages
 * +-----+
 * |   Mapped          |
 * ...
 * |   Unmapped        |
 * +-----+ XVM_MAX_BUFFER_SIZE * 512 + vpnoffset
```


XvMotion Buffer

- Per Migration buffer size of 64 MB
- Max of 8 = 512 MB
- Page bit map
 - Tracks which pages have been allocated
 - $64 \text{ MB} / 4\text{K} = 16\text{k bits} = 1 \text{ page} = 4\text{KB} = 16 \text{ kbits}$
- Block bitmaps
 - Tracks which blocks have been allocated
 - A block is 512 byte
 - 8 blocks per page
 - $64 \text{ MB} = 16\text{k pages} * 8 = 128\text{k blocks}$
 - 4 pages are required to track all blocks

XvMotion Buffer

- Bloom filter bitmap
 - One bloom filter per disk
 - We support 4 simultaneous disk copy
 - 2 pages for bloom filter
- Disk queue
 - 4 disk queues based on 4 simultaneous disk copy
 - A queue element is 16 bytes.
 - To have a queue of 128 k blocks we need 512 pages
- Buffer overhead
 - $1 + 4 + 2 + 512 = 521$ pages
 - $521 * 8 = 4168$ (~4k/128k) blocks lost in overhead!

XvMotion Buffer – Module init

- No heap management code
 - Neither static or nor dynamic heap
- We do our own buffer management
- Very efficient allocation!

init_module

XVMotion_Module

XVMotion_SetupVASpace

Migrate_ModuleAlloc (Slices)

VASpace_Reserve (8 * 64 MB)

Initialize the first slice i.e default

XvMotionBitmapInit (page, block and bloomfilter)

XvMotionQueueInit

XvMotion Buffer – Migration Init

- Reserve slice at module init time
 - We only allocate bitmaps and queues
 - Block pages are NOT allocated!
 - No guarantee that a XvMotion can succeed in low memory conditions

VMotion_PreCopyStart

VMotionPreCopyStart

XVMotion_SetupMigration

XVMotionAllocPool

XVMotionAllocSliceBitmap

For (numPages = metadata of 521 pages)

Migrate_NiceAllocPage

VASpace_Map

XVMotion Grow Buffer

- Block Pages
 - Even at Migration initialization time we don't allocate block pages
 - Allocated on demand

MigrateBridge_XVMAllocBlocks

 XVMotion_AllocBlocks

 XVMotionCheckForBlockMemLocked

 XVMotionGrowBuffer

 for (slice->pageBitmap.len)

 Migrate_NiceAllocPage

 PagePool_Alloc

 VASpace_Map

 If (freeblocks > reqblocks)

 break;

Block Allocation

- Source – Mirror IO

SVMAsyncIORemoteInt

MigrateBridge_XVMAllocBlocks

XVMotion_AllocBlocks

- Source - Clone - SVMAsyncIORead

SVMAsyncIORead

MigrateBridge_XVMAllocBlocks

XVMotion_AllocBlocks

- Destination - Stream

XVMotion_ReadPrepareBlocks

XVMotionCheckForBlockMemLocked

XvMotion Storage Stream

Storage Stream

- Stream is divided into generator and handler
 - Generator – Sender side
 - Handler – Receiver side
- Sender Side
 - Generator - XVMotion_GetStorageBlocks
 - Payload write - XVMotion_WriteBlocks
- Receiver Side
 - Handler - XVMotion_HandleBlocks
 - Read prepare - XVMotion_ReadPrepareBlocks
 - Read complete - XVMotion_ReadCompleteBlocks

Storage generator

- Generation - XVMotion_GetStorageBlocks
 - Prepare XVMotionBlockGroupData
 - Similar to VMotionPageGroupPublic
 - Same 128 block (vs Pages)
 - Round robin queue drain
 - Select the queue
 - Pick min blocks to transmit
 - MIN (readyForDrain, queuedForDraining)
 - Dequeue blocks from queue
 - Add address of the block to private data -> blockPtr
 - Maintain count of number of blocks

Storage write

- Payload write - XVMotion_WriteBlocks
- Transfer memory contents from XvMotion buffer to stream buffer
 - For (publicData.numblocks)
 - Copy the block data to completion data
 - Note: Completion buffers are allocated at channel init time
 - So no stream allocation

Handle a storage stream

- Handler - XVMotion_HandleBlocks
- Prepare segments from public data
 - For (public data => numblocks)
 - Determine total blocks
 - Determine total length of blocks
 - Populate the segment with total len
 - Completion world will pick the segment and process

Read storage stream

- Read prepare - XVMotion_ReadPrepareBlocks
 - Prepare a list of blocks to read from storage stream to XvMotion buffer
 - Get number of blocks from stream's public data num blocks.
 - Allocate num storage blocks from XvMotion buffer
 - XVMotionCheckForBlockMemLocked
 - XVMotionGrowBuffer (if required)
 - XVMotionAllocBlock
- Read complete - XVMotion_ReadCompleteBlocks
 - Read data from storage stream to XvMotion buffers prepared
 - For (pubGroup => numblocks)
 - Initialize a SGA
 - Populate SGA with pubGroup block data
 - XVMotionAddDiskIO

XvMotion IO

- Destination
- From storage stream add a new IO disk queue
 - XVMotionAddDiskIO
 - Allocate ioentry and make it point to SGA
 - XVMotionWaitForIOCount
 - `dq->ioCount < XVM_MAX_DST_OIO_COUNT`
 - XVMotionHandleConflictingIO
 - Check disk queue active IO
 - Is new IO to the same region as existing IO
 - If so enqueue and return
 - XVMotionAsyncIOWrite
 - `Async_PushCallbackFrameSafe`
 - `FSS_AsyncFileIO`

Thank You